



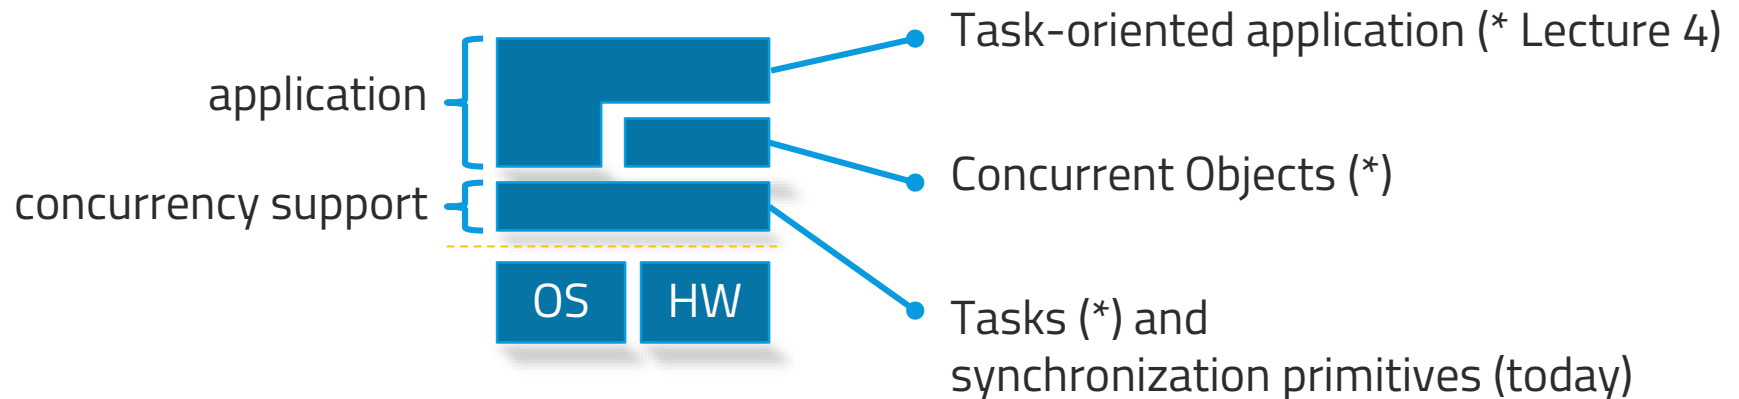
CSCI-GA.3033-017  
**Special Topics:**  
**Multicore Programming**

**Lecture 5**  
**Threads and Pthreads II**

Christopher Mitchell, Ph.D.  
cmitchell@cs.nyu.edu || <http://z80.me>

# Context

- We're exploring the layers of an application running on top of multicore hardware.



Mutices

Semaphores

Cond Vars

# Today's Agenda

- Add the grader; email both of us your commit hashes
  - Github: Nitesh (Github account: nits3392)
  - Email: ns3664@nyu.edu
- Concrete Pthread Creation
- A Simple Mutual Exclusion Algorithm
- Semaphores and “Blocking” Mutices
- Condition Variables and the Wakeup Problem



# Concrete Pthread Creation



```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 4

void *printHello(void *thread_id) {
    size_t tid = (size_t)thread_id;
    printf("Hello World! It's me, thread #%ld!\n", tid);
    pthread_exit(NULL);
}

int main(int argc, char *argv[]) {
    pthread_t threads[NUM_THREADS];
    for(size_t t = 0; t < NUM_THREADS; t++) {
        printf("In main: creating thread %ld\n", t);
        int rval = 0;
        if (0 != (rval = pthread_create(&threads[t], NULL,
            printHello, (void *)t))) {
            fprintf(stderr, "ERROR; return code from pthread_create() is %d\n", rval);
            break;
        }
        pthread_detach(threads[t]);
    }


    pthread_exit(NULL);
    return 0;
}
```

# Understanding Mutices



- The literature on mutual exclusion algorithms is extensive
- In practice, mutual exclusion is implemented with some form of hardware and OS support (coming soon!)
- But we'll first develop our algorithmic and correctness evaluation skills by looking at some very elegant solutions that use only shared memory

# Understanding Mutices: Review



- Mutual exclusion algorithms will provide two methods, `lock()` and `unlock()`, that allow us to mark the beginning and the end of critical sections in our code.
  - `lock()` will block the caller until there is no other thread in the critical section
  - `unlock()` would allow other thread in the critical section if there's one waiting

# Peterson's Algorithm

- Mutual exclusion solution in “three” lines of code.

```
int  victim;  
bool flag[2] = { false, false };
```

Current State

```
void lock(void) {  
    int me = my_tid % 2  
    int other = 1 - me  
    ① flag[me] = true  
    ② victim = me  
    ③ while (flag[other] &&  
           victim == me) {  
        no-op;  
    }  
}
```

```
void unlock(void) {  
    int me = my_tid % 2  
    flag[me] = false  
}
```

# Uncontested Case

- Solved in two simple tests and no wait.

```
int  victim = 0;  
bool flag[2] = { true, false };
```

Current State

```
void lock(void) {  
    int me = my_tid % 2  
    int other = 1 - me  
    flag[me] = true  
    victim = me  
    while (flag[other] &&  
           victim == me) {  
        no-op;  
    }  
}
```

Thread 0

```
void unlock(void) {  
    int me = my_tid % 2  
    flag[me] = false  
}
```



# Contested Case

- Thread 1 will busy-wait until Thread 0 executes unlock

```
int  victim = 1;  
bool flag[2] = { true, true };
```

Current State

```
void lock(void) {  
    int me = my_tid % 2  
    int other = 1 - me  
    flag[me] = true  
    victim = me  
    while (flag[other] &&  
           victim == me) {  
        no-op;  
    }  
}
```

Thread 1

loop

```
void unlock(void) {  
    int me = my_tid % 2  
    flag[me] = false  
}
```

Thread 0

Mutices

Semaphores

Cond Vars

# Simultaneous Case

- Assumes both cannot write to same mem addr at once.

```
int  victim = ?;  
bool flag[2] = { true, true };
```

Current State

```
void lock(void) {  
    int me = my_tid % 2  
    int other = 1 - me  
    flag[me] = true  
    victim = me  
    while (flag[other] &&  
           victim == me) {  
        no-op;  
    }  
}
```

```
void unlock(void) {  
    int me = my_tid % 2  
    flag[me] = false  
}
```

# Analyzing Peterson's Algorithm



- Three criteria
  1. Mutual exclusion
  2. Progress (no deadlock)
  3. Bounded waiting (fairness)

# Analyzing Peterson's Algorithm



- Does Peterson's Algorithm guarantee mutual exclusion?
- Assume that both threads could pass the tests:  
`while (flag[other] && victim == me)`
- This would have meant that each would have set the victim to be itself and each would have seen the victim as the other thread
- Is this possible?

# Analyzing Peterson's Algorithm



- Is it starvation free? Fair?
- If a thread `unlock()`s then it would set victim to be itself
- The contenting thread would have the chance to pass `lock()` then.
- Now, nothing prevents the `unlock()` thread from going ahead and trying a `lock()` again.
  - But in that case, that thread changes the victim to be itself before getting in the loop.



# Deadlock on Simultaneous Case

- Using flag[] alone doesn't work
- Could it be simpler? Instead of flag[] and victim, let's have one or the other

```
bool flag[2] = { true, true };
```

Current State

```
void lock(void) {  
    int me = my_tid % 2  
    int other = 1 - me  
    flag[me] = true  
    while (flag[other]) {  
        no-op;  
    }  
}
```

```
void unlock(void) {  
    int me = my_tid % 2  
    flag[me] = false  
}
```

# Deadlock on the Uncontested Case

- Peterson's is as simple as it gets.
- Now with 'victim' only

`int victim;`

Current State


```
void lock(void) {  
    int me = my_tid % 2  
    int other = 1 - me  
    victim = true  
    while (flag[other]) {  
        no-op;  
    }  
}
```

`void unlock(void)`

# Observations


- We'll be using the reasoning we developed here in other algorithms.
- Peterson's lock is important because it was arguably the first one to show 2-thread mutual exclusion can be solved in a **simple** way
- It makes several assumptions, though, and these turn out not to be so in practice
  - Number of threads is known *a priori*
  - Program execution is **strictly sequential** and each instruction is atomic
- The algorithm uses busy waiting

# Semaphores



- Invented by Edsger Dijkstra in 1968.
- A synchronization primitive that enables waiting without busy-wait
- A semaphore has an initial value, usually 1 (binary semaphore), and two operations: `up()` and `down()`
- There are some differences on that interface in practice (and some flavors of interface)
  - We'll reason using `up/down`

# Semaphores



- `down()`
  - Atomically check that value is greater than 0 and decrement it, allowing the thread to continue
  - Otherwise, suspend the thread, waiting on the counter value to be greater than 0
- `up()`
  - Atomically increment the counter.



# Exercise: Semaphore-Based Mutex



- Required methods: `lock()` and `unlock()`
- Assume unlocking an unlocked mutex is okay
  - ie, handle this case properly
  - Without needing to handle this, we can use a single semaphore initially set to 1 as a mutex!
- Semaphores are powerful!

# Semaphore-Based "Blocking" Mutex




```
S: private semaphore, initial count 1
locked: Boolean, initially false
holder: thread ID
M: private semaphore, initial count 1
```

```
void lock(void) {
    M.down();
    S.down(); // Critical v
    locked = true;
    holder = self();
    S.up();    // Critical ^
}
```

```
void unlock(void) {
    S.down(); // Critical v
    if (!locked ||
        holder != self()) {
        S.up();
        return;
    }
    locked = false;
    S.up();    // Critical ^
    M.up();
}
```

# Semaphore-Based "Blocking" Mutex



- Starvation free!
  - `lock()` can block at `M.down()`, but only until `unlock()`
  - Nothing can prevent `unlock()`

```
void lock(void) {  
    M.down();  
    S.down(); // Critical v  
    locked = true;  
    holder = self();  
    S.up();   // Critical ^  
}
```

```
void unlock(void) {  
    S.down(); // Critical v  
    if (!locked ||  
        holder != self()) {  
        S.up();  
        return;  
    }  
    locked = false;  
    S.up();   // Critical ^  
    M.up();  
}
```

# Semaphore-Based "Blocking" Mutex

- Deadlock free!
  - S protects the critical sections, never left locked.
  - No cycles

```
void lock(void) {  
    M.down();  
    S.down(); // Critical v  
    locked = true;  
    holder = self();  
    S.up();    // Critical ^  
}
```

```
void unlock(void) {  
    S.down(); // Critical v  
    if (!locked ||  
        holder != self()) {  
        S.up();  
        return;  
    }  
    locked = false;  
    S.up();    // Critical ^  
    M.up();  
}
```


# Semaphore Considerations



- In order to sleep instead of busy wait, need to call the OS
  - But if at each semaphore operation we incur one system call, doing synchronization may be expensive
  - “Ideal” mutex (“blocking”)
    - Use shared memory to store the state of the lock
    - Uncontended case requires access to memory only
    - In contended case, ask OS to sleep until value of shared memory changes
- Reasoning about semaphore algorithms can be daunting.



# Condition Variable Semantics



- Allows a thread to wait on a given predicate to change
- Associated with a mutex that protects the predicate state
- Operations
  - `wait()` – atomatically suspends the execution of the thread and unlock the associated mutex
  - `signal()` – if there's at least one thread suspended on the cond var, then dequeue it and resume execution, again, atomically
  - `broadcast()` – if there are any threads suspended on the cond var, resume execution for all of them. They'll contend for the associated lock.

# Implementing wait() and signal()

- `wait()` is always inside a loop that checks the predicate
  - Easier to implement in terms of thread scheduling
  - Allows signal on every predicate change
- `signal()` should be done inside the lock (although it might be correct to do so outside)

```
element dequeue()  
    pthread_mutex_lock(queue_lock)  
    ...  
    while (empty) {  
        pthread_cond_wait(cond, queue_lock)  
    }  
    // !empty  
    ...
```

```
enqueue(element)  
    pthread_mutex_lock(queue_lock)  
    ...  
    empty = false  
    pthread_cond_signal(cond)  
    ...
```

# Fair Condition Variables Using Semaphores

- Reported by Andrew Birrell in 1993<sup>1</sup>

cond has: a queue of waiters and a semaphore  
thread has: a private semaphore with initial count 0

```
wait(cond, mutex) {  
    // assumes mutex held  
    // assumes cond is false  
    Thread self = Thread.self()  
    cond.sem.down()  
    enqueue(self) on cond's queue  
    cond.sem.up()  
    release mutex  
    self.sem.down()  
    acquire mutex  
}
```

```
signal(cond) {  
    // assumes mutex held  
    cond.sem.down()  
    if a thread t is waiting {  
        dequeue(t)  
        t.sem.up()  
    }  
    cond.sem.up()  
}
```

<sup>1</sup><https://birrell.org/andrew/papers/ImplementingCVs.pdf>

# Fair Condition Variables Using Semaphores

- Reported by Andrew Birrell in 1993

cond has: a queue of waiters and a semaphore  
thread has: a private semaphore with initial count 0

```
wait(cond, mutex) {  
    // assumes mutex held  
    // assumes cond is false  
    Thread self = Thread.self()  
    cond.sem.down()  
    enqueue(self) on cond's queue  
    cond.sem.up()  
    release mutex  
    self.sem.down()  
    acquire mutex  
}
```

```
signal(cond) {  
    // assumes mutex held  
    cond.sem.down()  
    if a thread t is waiting {  
        dequeue(t)  
        t.sem.up()  
    }  
    cond.sem.up()  
}
```

Equivalent to atomically releasing the lock and going to sleep. What if `t.sem.up()` runs before `self.sem.down()`? What could happen if we released the mutex lock earlier?

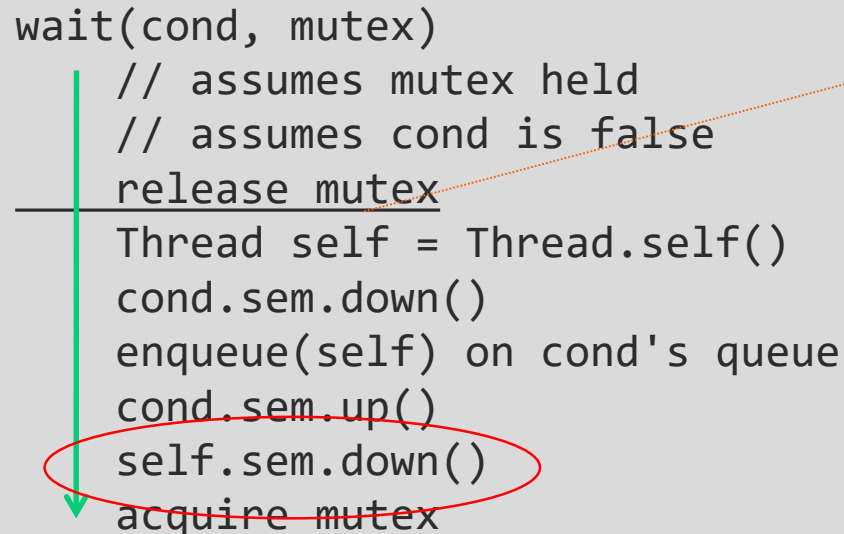
# Lost Wakeup Race Condition

- A deadly (and common) mistake

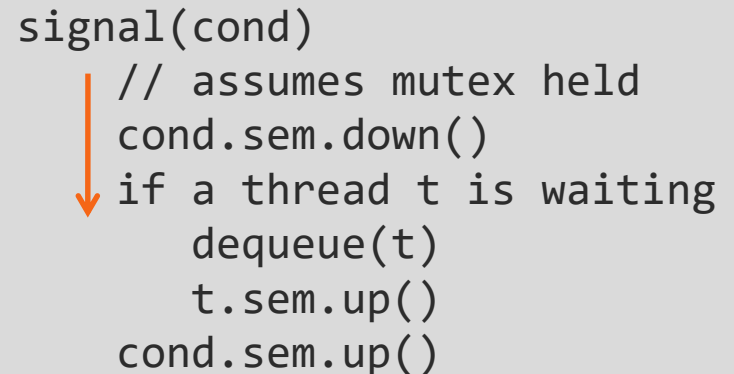
```
wait(cond, mutex)
// assumes mutex held
// assumes cond is false
release mutex


---


Thread self = Thread.self()
cond.sem.down()
enqueue(self) on cond's queue
cond.sem.up()
self.sem.down()
acquire mutex
```

A diagram illustrating the logic of the wait function. A green vertical arrow on the left points downwards, indicating the flow of execution. A red oval highlights the 'self.sem.down()' line, which is annotated with 'acquire mutex' below it. A red dotted line points from the 'release mutex' line to the 'self.sem.down()' line, highlighting a race condition where the mutex is released before the thread has fully acquired it.

```
signal(cond)
// assumes mutex held
cond.sem.down()
↓ if a thread t is waiting
  dequeue(t)
  t.sem.up()
cond.sem.up()
```

A diagram illustrating the logic of the signal function. An orange arrow points downwards from the 'cond.sem.down()' line to the 'if a thread t is waiting' block, indicating the flow of execution.



# Conclusion



- Mutices: How can we build our own?
- Semaphores: Making mutices, working with the OS
- Cond Vars: Implementation gotchas
- Lab 1: Questions? Problems?