



CSCI-GA.3033-017
Special Topics:
Multicore Programming

Lecture 6
Coordinating Resources

Christopher Mitchell, Ph.D.
cmitchell@cs.nyu.edu || <http://z80.me>


Homework 1 Review

Types of Parallelism



- Instruction-Level Parallelism [CPU]
 - Pipelining, requires pipelined CPU
- Basic Block Parallelism [CPU] [Compiler]
 - Reordering and parallelizing instructions within a block
 - Parallelizing instructions from multiple blocks
 - Requires register copies and functional unit copies
- Loop Level Parallelism [Compiler]
 - Interleave and parallelize instructions from multiple iterations
- Task Parallelism [Programmer]
 - Threads: related work, often sharing same memory space
- Process Parallelism [Programmer]
 - Distinct work to be completed in parallel
- Machine Parallelism [Programmer]
 - Break work into groups of related processes spread across multiple machines


Outline



- ~~Homework Review~~
- Coordinating Resources: Reasoning about two mutex/semaphore-based schemes
 - Reader-Writer Locks
 - Barriers
- Lab 2 Techniques
 - Socket Refresher
 - Thread Pools

Outline

- ~~Homework Review~~
- Coordinating Resources
 - [Reader-Writer Locks](#)
 - Barriers
- Lab 2 Techniques
 - Socket Refresher
 - Thread Pools



The Reader-Writer Problem




- Consider a resource
 - Shared by several threads
 - Some threads may only want to read
 - Others may want to modify
- Could we coordinate these writers and readers?
- Idea: a reader-writer lock [pair]
 - Each reader acquires a special lock that allows them to share the resource with other readers
 - A writer acquires another kind of lock that gives it exclusive access to the resource
 - The locks work in tandem to guarantee the resource's consistency

POSIX File Reader-Writer Lock

- File locking between processes or threads
- `flock(file_handle, mode)`
 - `LOCK_SH`: Shared (reader) lock
 - `LOCK_EX`: Exclusive (writer) lock
 - Bitwise OR with `LOCK_NB`: Nonblocking

Process 1:

```
FILE* fh = fopen(F);  
flock(fh, LOCK_SH);
```




POSIX File Reader-Writer Lock

- File locking between processes or threads
- `flock(file_handle, mode)`
 - `LOCK_SH`: Shared (reader) lock
 - `LOCK_EX`: Exclusive (writer) lock
 - Bitwise OR with `LOCK_NB`: Nonblocking

Process 2:

```
FILE* fh = fopen(F);  
flock(fh, LOCK_SH);
```




POSIX File Reader-Writer Lock

- File locking between processes or threads
- `flock(file_handle, mode)`
 - `LOCK_SH`: Shared (reader) lock
 - `LOCK_EX`: Exclusive (writer) lock
 - Bitwise OR with `LOCK_NB`: Nonblocking

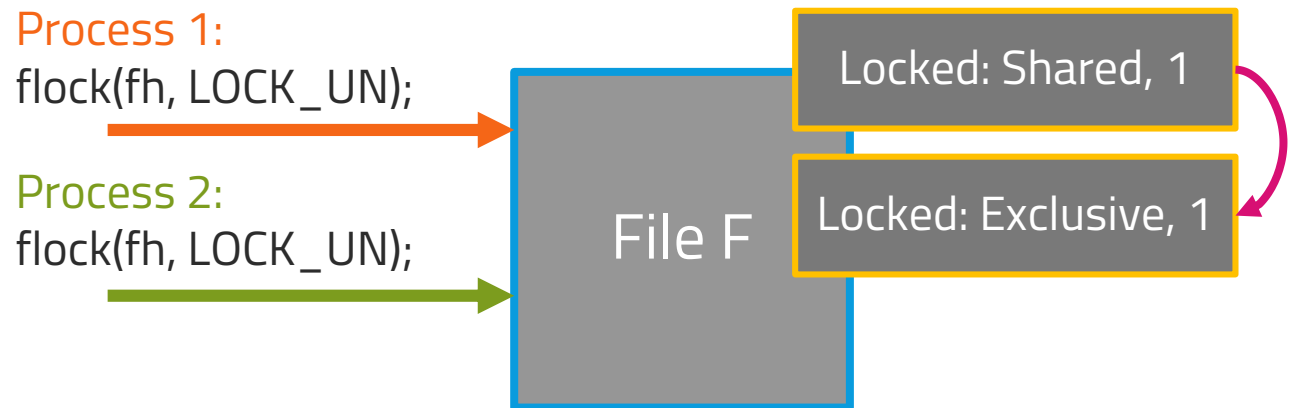
Process 3:

```
FILE* fh = fopen(F);  
flock(fh, LOCK_EX);
```



POSIX File Reader-Writer Lock

- File locking between processes or threads
- `flock(file_handle, mode)`
 - LOCK_SH: Shared (reader) lock
 - LOCK_EX: Exclusive (writer) lock
 - Bitwise OR with LOCK_NB: Nonblocking



Simple Reader-Writer Lock



- Forgot files: let's implement a simple reader-writer lock
- Semantics:
 - Allow any number of shared readers
 - Allow a single exclusive writer
 - Fairness? Worry about it later
- Toolset
 - Mutices

Simple Reader-Writer Lock

```
int read_count = 0
mutex mut_read, write_lock
```

```
reader_lock():
    lock(mut_read)
    read_count += 1
    if read_count == 1:
        lock(write_lock)
    unlock(mut_read)
```

```
reader_unlock():
    lock(mut_read)
    read_count -= 1
    if read_count == 0:
        unlock(write_lock)
    unlock(mut_read)
```

```
writer_lock():
    lock(write_lock)

writer_unlock():
    unlock(write_lock)
```

Who gets the priority? Readers or writers?

Simple Reader-Writer Lock



Reader arrives before writer

```
int read_count = 1
mutex mut_read, write_lock
```

```
reader_lock():
    lock(mut_read)
    read_count += 1
    if read_count == 1:
        lock(write_lock)
    unlock(mut_read)

reader_unlock():
    lock(mut_read)
    read_count -= 1
    if read_count == 0:
        unlock(write_lock)
    unlock(mut_read)
```

```
writer_lock():
    lock(write_lock)

writer_unlock():
    unlock(write_lock)
```

Simple Reader-Writer Lock

Reader arrives before writer

```
int read_count = 1
mutex mut_read, write_lock
```

```
reader_lock():
    lock(mut_read)
    read_count += 1
    if read_count == 1:
        lock(write_lock)
    unlock(mut_read)

reader_unlock():
    lock(mut_read)
    read_count -= 1
    if read_count == 0:
        unlock(write_lock)
    unlock(mut_read)
```

```
writer_lock():
    lock(write_lock)

writer_unlock():
    unlock(write_lock)
```


Simple Reader-Writer Lock: Starvation



Second reader arrives before first reader finishes

```
int read_count = 2
mutex mut_read, write_lock
```

```
reader_lock():
    lock(mut_read)
    read_count += 1
    if read_count == 1:
        lock(write_lock)
    unlock(mut_read)
```

```
reader_unlock():
    lock(mut_read)
    read_count -= 1
    if read_count == 0:
        unlock(write_lock)
    unlock(mut_read)
```

```
writer_lock():
    lock(write_lock)
writer_unlock():
    unlock(write_lock)
```

Reader-Writer Lock v2



Give writers priority over readers.

```
int read_count, write_count
mutex mut_read, mut_write, read_lock, write_lock
```

```
reader_lock():
    lock(read_lock)
    lock(mut_read)
    read_count += 1
    if read_count == 1:
        lock(write_lock)
    unlock(mut_read)
    unlock(read_lock)

reader_unlock():
    lock(mut_read)
    read_count -= 1
    if read_count == 0:
        unlock(write_lock)
    unlock(mut_read)
```

```
writer_lock():
    lock(mut_write)
    write_count += 1
    if write_count == 1:
        lock(read_lock)
    unlock(mut_write)
    lock(write_lock)

writer_unlock():
    lock(mut_write)
    write_count -= 1
    if write_count == 0:
        unlock(read_lock)
    unlock(mut_write)
    unlock(write_lock)
```

Reader-Writer Lock v2



One reader, then one writer, arrives.

```
int read_count = 1, write_count
mutex mut_read, mut_write, read_lock, write_lock
```

```
reader_lock():
    lock(read_lock)
    lock(mut_read)
    read_count += 1
    if read_count == 1:
        lock(write_lock)
    unlock(mut_read)
    unlock(read_lock)
```

```
reader_unlock():
    lock(mut_read)
    read_count -= 1
    if read_count == 0:
        unlock(write_lock)
    unlock(mut_read)
```

```
writer_lock():
    lock(mut_write)
    write_count += 1
    if write_count == 1:
        lock(read_lock)
    unlock(mut_write)
    lock(write_lock)
```

```
writer_unlock():
    lock(mut_write)
    write_count -= 1
    if write_count == 0:
        unlock(read_lock)
    unlock(mut_write)
    unlock(write_lock)
```

Reader-Writer Lock v2



Second reader arrives.

```
int read_count = 0, write_count = 2
mutex mut_read, mut_write, read_lock, write_lock
```

```
reader_lock():
    ↓ lock(read_lock)
    lock(mut_read)
    read_count += 1
    if read_count == 1:
        lock(write_lock)
    unlock(mut_read)
    unlock(read_lock)

reader_unlock():
    lock(mut_read)
    read_count -= 1
    if read_count == 0:
        unlock(write_lock)
    unlock(mut_read)
```

```
writer_lock():
    lock(mut_write)
    write_count += 1
    if write_count == 1:
        lock(read_lock)
    unlock(mut_write)
    ↓ lock(write_lock)

writer_unlock():
    lock(mut_write)
    write_count -= 1
    if write_count == 0:
        unlock(read_lock)
    unlock(mut_write)
    unlock(write_lock)
```

Reader-Writer Lock v2



Now writers can starve readers.

```
int read_count = 0, write_count = 2
mutex mut_read, mut_write, read_lock, write_lock
```

```
reader_lock():
    ↓ lock(read_lock)
    lock(mut_read)
    read_count += 1
    if read_count == 1:
        lock(write_lock)
    unlock(mut_read)
    unlock(read_lock)

reader_unlock():
    lock(mut_read)
    read_count -= 1
    if read_count == 0:
        unlock(write_lock)
    unlock(mut_read)
```

```
writer_lock():
    lock(mut_write)
    write_count += 1
    if write_count == 1:
        lock(read_lock)
    unlock(mut_write)
    lock(write_lock)

writer_unlock():
    lock(mut_write)
    write_count -= 1
    if write_count == 0:
        unlock(read_lock)
    unlock(mut_write)
    unlock(write_lock)
```


Reader-Writer Lock v3



Tracing a reader, then a writer

```
int a_readers, a_writers, p_readers, p_writers    // Active & pending
mutex mut, cond_var read_cond, write_cond
```

```
reader_lock():
    lock(mut)
    while a_writers + p_writers:
        p_readers += 1
        read_cond.wait(mut)
        p_readers -= 1
    a_readers += 1
    unlock(mut)
```

```
reader_unlock():
    lock(mut)
    a_readers -= 1
    if !a_readers && p_writers:
        write_cond.signal()
    unlock(mut)
```

```
writer_lock():
    lock(mut)
    while a_writers + a_readers:
        p_writers += 1
        write_cond.wait(mut)
        p_writers -= 1
    a_writers += 1
    unlock(mut)
```

```
writer_unlock():
    lock(mut)
    a_writers -= 1
    if p_writers:
        write_cond.signal()
    else if p_readers:
        read_cond.broadcast()
    unlock(mut)
```

Reader-Writer Lock v3



Tracing a reader, then a writer

```
int a_readers, a_writers, p_readers, p_writers    // Active & pending
mutex mut, cond_var read_cond, write_cond
```

```
reader_lock():
    lock(mut)
    while a_writers + p_writers:
        p_readers += 1
        read_cond.wait(mut)
        p_readers -= 1
    a_readers += 1
    unlock(mut)
```

```
reader_unlock():
    lock(mut)
    a_readers -= 1
    if !a_readers && p_writers:
        write_cond.signal()
    unlock(mut)
```

```
writer_lock():
    lock(mut)
    while a_writers + a_readers:
        p_writers += 1
        write_cond.wait(mut)
        p_writers -= 1
    a_writers += 1
    unlock(mut)
```

```
writer_unlock():
    lock(mut)
    a_writers -= 1
    if p_writers:
        write_cond.signal()
    else if p_readers:
        read_cond.broadcast()
    unlock(mut)
```

Reader-Writer Lock v3



Tracing a reader, a writer, and a second writer.

```
int a_readers, a_writers, p_readers, p_writers    // Active & pending
mutex mut, cond_var read_cond, write_cond
```

```
reader_lock():
    lock(mut)
    while a_writers + p_writers:
        p_readers += 1
        read_cond.wait(mut)
        p_readers -= 1
    a_readers += 1
    unlock(mut)
```

```
reader_unlock():
    lock(mut)
    a_readers -= 1
    if !a_readers && p_writers:
        write_cond.signal()
    unlock(mut)
```

```
writer_lock():
    lock(mut)
    while a_writers + a_readers:
        p_writers += 1
        write_cond.wait(mut)
        p_writers -= 1
    a_writers += 1
    unlock(mut)
```

```
writer_unlock():
    lock(mut)
    a_writers -= 1
    if p_writers:
        write_cond.signal()
    else if p_readers:
        read_cond.broadcast()
    unlock(mut)
```

Reader-Writer Lock v3

Tracing a reader, a writer, and a second writer.

```
int a_readers, a_writers, p_readers, p_writers    // Active & pending
mutex mut, cond_var read_cond, write_cond
```

```
reader_lock():
    lock(mut)
    while a_writers + p_writers:
        p_readers += 1
        read_cond.wait(mut)
        p_readers -= 1
    a_readers += 1
    unlock(mut)
```


```
reader_unlock():
    lock(mut)
    a_readers -= 1
    if !a_readers && p_writers:
        write_cond.signal()
    unlock(mut)
```

```
writer_lock():
    lock(mut)
    while a_writers + a_readers:
        p_writers += 1
        write_cond.wait(mut)
        p_writers -= 1
    a_writers += 1
    unlock(mut)
```

```
writer_unlock():
    lock(mut)
    a_writers -= 1
    if p_writers:
        write_cond.signal()
    else if p_readers:
        read_cond.broadcast()
    unlock(mut)
```

Choose
priority here

Final Reader-Writer Lock Miscellanea



- Every time we see a structure taking many readers, R/W seem the thing to do.
- However...
 - Even in the reader-only case, there could be contention on the reader counter mutex.
 - Maintaining fairness can cause contention
- Recent work:
 - "Scalable Reader-Writer Locks", from Lev et al. 2009


Pthread Reader-Writer Lock



- Type: `pthread_rwlock_t`
- Initialization: `int pthread_rwlock_init(pthread_rwlock_t *rwlock, const pthread_rwlockattr_t *attr);`
- Lock for read:
 - Blocking:
`int pthread_rwlock_rdlock(pthread_rwlock_t *rwlock);`
 - Nonblocking:
`int pthread_rwlock_tryrdlock(pthread_rwlock_t *rwlock);`
- Lock for write
 - Blocking:
`int pthread_rwlock_wrlock(pthread_rwlock_t *rwlock);`
 - Nonblocking:
`int pthread_rwlock_trywrlock(pthread_rwlock_t *rwlock);`

Outline

- ~~Homework Review~~
- Coordinating Resources
 - Reader-Writer Locks
 - **Barriers**
- Lab 2 Techniques
 - Socket Refresher
 - Thread Pools




Barrier



- Synchronize group of threads at single point
 - Each thread waits until all threads arrive
 - Each thread continues
- Solution
 - Mutex or semaphore to count arrivals
 - Mutex or semaphore to hold threads until count is equal to number of threads

Simple Semaphore-Based Barrier




```
semaphore arrival = 1, departure = 0;
int counter = 0, int n = num_threads;

void await(void) {
    arrival.down();    // Acts as mutex & block on arrival
    counter += 1;
    if (counter < n) {
        arrival.up();
    } else {
        departure.up();
    }
    departure.down(); // Acts as mutex & block on departure
    counter -= 1;
    if (counter > 0) {
        departure.up();
    } else {
        arrival.up(); // Back to initial conditions
    }
}
```

Must be known a priori

Simple Semaphore-Based Barrier




First arrival

```
semaphore arrival = 1, departure = 0;
int counter = 1, int n = num_threads;

void await(void) {
    arrival.down();    // Acts as mutex & block on arrival
    counter += 1;
    if (counter < n) {
        arrival.up();
    } else {
        departure.up();
    }
    departure.down(); // Acts as mutex & block on departure
    counter -= 1;
    if (counter > 0) {
        departure.up();
    } else {
        arrival.up(); // Back to initial conditions
    }
}
```


Simple Semaphore-Based Barrier




$n - 1$ arrivals

```
semaphore arrival = 1, departure = 0;
int counter = n - 1, int n = num_threads;

void await(void) {
    arrival.down();    // Acts as mutex & block on arrival
    counter += 1;
    if (counter < n) {
        arrival.up();
    } else {
        departure.up();
    }
    departure.down(); // Acts as mutex & block on departure
    counter -= 1;
    if (counter > 0) {
        departure.up();
    } else {
        arrival.up(); // Back to initial conditions
    }
}
```

Simple Semaphore-Based Barrier




n arrivals

```
semaphore arrival = 0, departure = 1;
int counter = 0, int n = num_threads;

void await(void) {
    arrival.down();    // Acts as mutex & block on arrival
    counter += 1;
    if (counter < n) {
        arrival.up();
    } else {
        departure.up();
    }
    departure.down(); // Acts as mutex & block on departure
    counter -= 1;
    if (counter > 0) {
        departure.up();
    } else {
        arrival.up(); // Back to initial conditions
    }
}
```

Simple Semaphore-Based Barrier




n arrivals, 1 departure

```
semaphore arrival = 0, departure = 1;
int counter = n - 1, int n = num_threads;

void await(void) {
    arrival.down();    // Acts as mutex & block on arrival
    counter += 1;
    if (counter < n) {
        arrival.up();
    } else {
        departure.up();
    }
    departure.down(); // Acts as mutex & block on departure
    counter -= 1;
    if (counter > 0) {
        departure.up();
    } else {
        arrival.up(); // Back to initial conditions
    }
}
```

Simple Semaphore-Based Barrier




n arrivals, n - 1 departures

```
semaphore arrival = 0, departure = 1;
int counter = 1, int n = num_threads;

void await(void) {
    arrival.down();    // Acts as mutex & block on arrival
    counter += 1;
    if (counter < n) {
        arrival.up();
    } else {
        departure.up();
    }
    departure.down(); // Acts as mutex & block on departure
    counter -= 1;
    if (counter > 0) {
        departure.up();
    } else {
        arrival.up(); // Back to initial conditions
    }
}
```

Simple Semaphore-Based Barrier




n arrivals

```
semaphore arrival = 1, departure = 0;
int counter = n - 1, int n = num_threads;

void await(void) {
    arrival.down();    // Acts as mutex & block on arrival
    counter += 1;
    if (counter < n) {
        arrival.up();
    } else {
        departure.up();
    }
    departure.down(); // Acts as mutex & block on departure
    counter -= 1;
    if (counter > 0) {
        departure.up();
    } else {
        arrival.up(); // Back to initial conditions
    }
}
```

Pthread Barrier




- Surprise! Pthread has a barrier primitive
- Type: `pthread_barrier_t`
- Initialization:

```
int pthread_barrier_init(pthread_barrier_t* barrier,  
attributes, unsigned int count);
```
- Wait:


```
int pthread_barrier_wait(pthread_barrier_t* barrier);
```

Outline




- ~~Homework Review~~
- Coordinating Resources
 - Reader-Writer Locks
 - Barriers
- Lab 2 Techniques
 - Socket Refresher
 - Thread Pools

Lab 2




- Make our concurrent key-value store more useful: a multi-threaded key-value store server
 1. Implement reader-writer lock(s)
 2. Implement thread pool
 3. Implement GET/POST/DELETE frontend
- Three weeks to complete
- Due November 13, 2017

Lab 2



- Make our concurrent key-value store more useful: a multi-threaded key-value store server
 1. Implement reader-writer lock(s) -> **easy** (pthreads!)
 2. Implement value hashing and storage -> **moderate**
 3. Implement thread pool -> **challenging**
 4. GET/POST/DELETE frontend -> **provided**
- Three weeks to complete
- Due November 13, 2017

Socket (Re-)Primer



- Review: http://www.linuxhowtos.org/C_C++/socket.htm
- Relevant: <http://www.linuxhowtos.org/data/6/server2.c>
 - Please don't copy it, but good reference
- Concepts:
 - Socket connection (TCP: connectionful)
 - Passive (`listen()`ing/`accept()`ing) side
 - Active (`connect()`ing) side
- Server:
 - `listen()`
 - Repeatedly `accept()` -> use fd -> close fd

GET/POST/DELETE

- HTTP 1.1: <https://tools.ietf.org/html/rfc2616> (ouch)
- Saner: <https://www.jmarshall.com/easy/http/#sample>

GET /path HTTP/1.1 header header <i>[blank line]</i>	POST /path HTTP/1.1 header Content-Length: XXXX <i>[blank line]</i> contents	DELETE /path HTTP/1.1 header header <i>[blank line]</i>
HTTP/1.1 200 OK Content-Length: XXXX <i>[blank line]</i> contents	HTTP/1.1 200 OK Content-Length: XXXX <i>[blank line]</i> contents	HTTP/1.1 200 OK Content-Length: XXXX <i>[blank line]</i> contents

Note: newline is \r\n; see <https://www.w3.org/Protocols/rfc2616/rfc2616-sec2.html#sec2.2>

Thread Pool



- Thread work can be small pieces
 - Creating and destroying threads is expensive
 - Reduce overhead: reuse threads
1. Create group of N threads
 2. Use thread-safe queue to identify “idle” threads
 3. Atomically remove and invoke an idle thread when new work arrives
 4. Atomically add self back to queue when work is done