

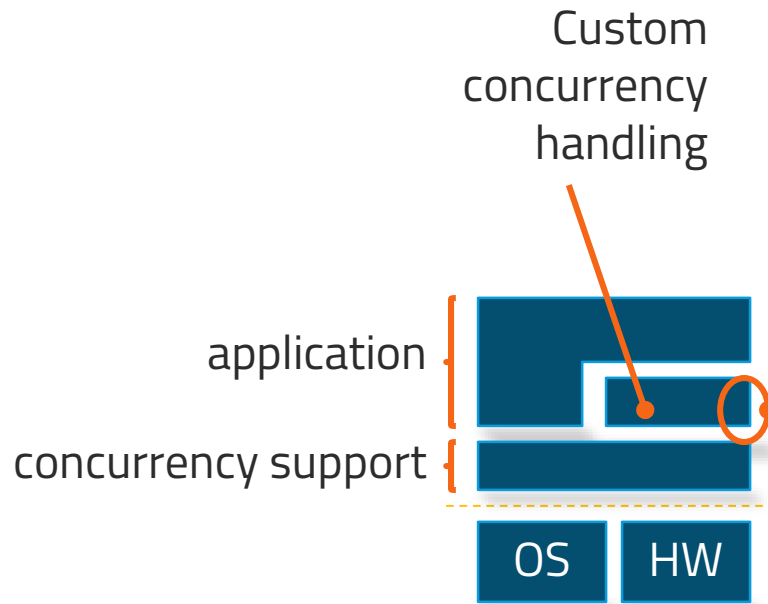


CSCI-GA.3033-017  
**Special Topics:**  
**Multicore Programming**

**Lecture 7**  
**Synchronized Structures Part 1**

Christopher Mitchell, Ph.D.  
cmitchell@cs.nyu.edu || <http://z80.me>

# Context




We're going to look at how to create more higher level synchronized data structures

- This week: Concurrent queues
- Next week: Building up to a lock-free hash table

# Outline

- Lock-Based Concurrent Queue
- Unlocking: A Lock-Free Concurrent Queue
- Understanding The ABA Problem
- Solving ABA: Memory Management and Reuse

# Lock-Based Concurrent Queue





- Consider a naïve concurrent queue
  - How many locks?
  - What problems with 1? (Correctness, performance, ...?)
  - What problems with 2? (Correctness, performance, ...?)
- Concept Reminder: Scoped locks
  - Automatically unlocked when destroyed
  - C++ scoped lock:  
`std::lock_guard<...> scoped_lock(my_mutex);`



# Concurrent Queue: Protecting Shared State


```
elem* dequeue() {  
    lock_guard(mutex);  
    elem* node = nullptr;  
    if (head != nullptr) {  
        node = head;  
        head = head->next;  
    }  
    if (head == nullptr) {  
        tail = nullptr;  
    }  
    return node;  
    // mutex unlocked  
}
```

```
enqueue(const& cnode) {  
    lock_guard(mutex);  
    elem* node =  
        new elem(cnode)  
    if (tail != nullptr) {  
        tail->next = node;  
    } else {  
        head = node;  
    }  
    tail = node;  
    // mutex unlocked  
}
```



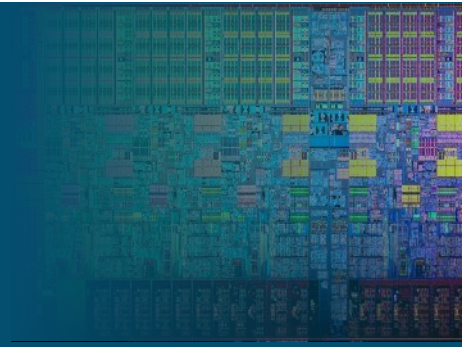
Both methods touch  
both head and tail.

# Concurrent Queue: Considering Correctness



- enqueue() and deque() must be protected under the same single lock
  - Why?
- Correctness: [Linearizability](#)
  - We dealt with this in Lab 1
  - Definition: There is *some* total serial order among all operations
    - Corollary: Each operation appears instantaneous
    - Corollary: No operation can see the intermediate state of another


# Concurrent Queue: Considering Correctness



- Can't split the lock as-is: no more linearizability
  - Another operation could see tail update before head update, or vice versa
- Linearization point: right after lock released
  - If there's no obvious linearization point, algorithm may be wrong, hard to reason about, or both.
- We can still split the lock and be linearizable!
  - How?

# Towards a Lock-Free Concurrent Queue


- Key concept: Sentinel
  - Dummy node, always first in the list (ie, what head points to)
  - Make enqueue() and dequeue() touch distinct portions of the state
  - dequeue() checks if tail and head point to same element
  - Assume pointer reads and writes are atomic operations
- Still a blocking, locked structure






# Dequeuing with Sentinel

- `dequeue()`: only touches head portion




# Enqueuing with Sentinel

- enqueue(): only touches tail portion



# Simultaneous Enqueue and Dequeue on Empty Queue

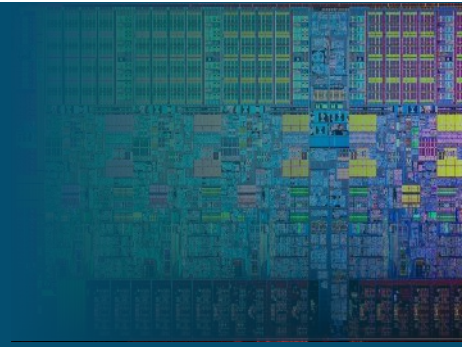
- Depends on ordering of tail pointer write in dequeue() and read in enqueue()
- Write->write reordering would break this! (Why?)





# A Lock-Free Queue

# Lock-Free Queue?



- It's possible!
- Current locks: 2
  - Single concurrent enqueue()
  - Single concurrent dequeue()
  - No protection between enqueue() and dequeue()
- Necessary atomic primitive: CAS
  - Compare and Swap



# Compare-And-Swap

CAS(address, expected, desired)

Atomically:

1. Check if (\*address == expected)
2. If so, set \*address = desired, return true
3. Otherwise, return false

- C++11:

std::atomic<T> type -> variable + locking!


template<class T>

bool atomic\_compare\_exchange\_strong(volatile [std::atomic](#)<T>\* obj,  
T\* expected, T desired);

# "One"-Step Lock-Free Enqueue


```
enqueue(elem* node) {  
    ...  
    tail->next = node;  
    CAS1{  
        if (tail->next == node)  
            tail = tail->next;  
    }  
}
```

Why can't we use this? ☹




# Two-Step Lock-Free Enqueue

```
enqueue(elem* node) {  
    ...  
    elem* cur_tail = tail;  
    CAS1{  
        if (tail->next == cur_tail->next)  
            tail->next = node;  
    }  
}
```




# Two-Step Lock-Free Enqueue

```
enqueue(elem* node) {  
    ...  
    elem* cur_tail = tail;  
    CAS1{  
        if (tail->next == cur_tail->next)  
            tail->next = node;  
    }  
    CAS2{  
        if (tail == cur_tail)  
            tail = node;  
    }  
}
```



# Two-Step Lock-Free Enqueue: Complete Code


```
enqueue(const value& val) {  
    elem* node = new elem();  
    node->val = val;  
    node->next = nullptr;  
    while(true) {  
        elem* cur_tail = tail;  
        if (CAS(tail->next, cur_tail->next,  
                node))  
        {  
            if (CAS(tail, cur_tail, node)) {  
                break;  
            }  
        }  
    }  
}
```





# Two-Step Lock-Free Enqueue Gotcha

- What happens during the two CAS operations?
- Tail might not point to last node
  - Why? On CAS2, node may not be real tail; another thread may have set tail->next to own node in CAS1.
- New invariant: tail may only point to last node or second-to-last node




# Two-Step Lock-Free Enqueue Gotcha


Thread 1 wants to enqueue node B  
Thread 2 wants to enqueue node C



1. Thread 1 caches tail in `cur_tail`, performs CAS 1.




2. Thread 2 caches tail in `cur_tail`, performs CAS 1




# Two-Step Lock-Free Enqueue Gotcha

Thread 1 wants to enqueue node B  
Thread 2 wants to enqueue node C



3. Thread 1 finishes enqueue with CAS 2. Tail hasn't moved, so its CAS 2 succeeds




# Two-Step Lock-Free Enqueue Gotcha

Thread 1 wants to enqueue node B  
Thread 2 wants to enqueue node C



4. Thread 2 tries to finish enqueue with CAS 2. It fails, because tail has changed. It starts over and performs CAS 1 and CAS 2.



5. If enqueue checked that tail->next was nullptr before updating tail->next, this problem would be avoided, but the tail pointer could then become stale.


# Two-Step Lock-Free Enqueue Fix

```
enqueue(const value& val) {
    elem* node = new elem();
    node->val = val;
    node->next = nullptr;
    while(true) {
        elem* cur_tail = tail;
        elem* next = cur_tail->next;
        if (cur_tail == tail) {
            if (next == nullptr) {
                if (CAS(tail->next, next, node)) {
                    break;
                }
            } else {
                CAS(tail, cur_tail, next); // Correct stale tail
            }
        }
    }
    // If the following fails, someone else will handle it.
    CAS(tail, cur_tail, node);
}
```



# Two-Step Lock-Free Dequeue

- What if dequeue() misses a node because of invariant (ie, tail points to second-to-last node)?
  - We'll get this on the next slide.
- General lock-free approach:



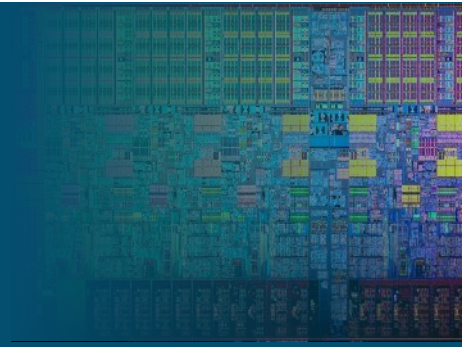
```
value dequeue() {  
    while(true) {  
        elem* cur_head = head;  
        elem* cur_next = head->next;  
        if (head == tail) {  
            return null_val;  
        }  
        value val = head->next;  
        if (CAS(head, cur_head, cur_next)) {  
            free(cur_head); // Old sentinel  
            return val;  
        }  
    }  
}
```

# Two-Step Lock-Free Dequeue

```
val dequeue() {
  value val = null_val;
  elem* free_node = nullptr;
  while(true) {
    elem* cur_head = head;
    elem* cur_tail = tail;
    elem* next = cur_head->next; // Catch half-done enqueue()!
    if (cur_head == cur_tail) {
      if (next == nullptr) {
        return nullptr; // Empty queue
      }
      CAS(tail, cur_tail, next); // Fix: tail->next != nullptr
    } else {
      free_node = head;
      val = next->val;
      if (CAS(head, cur_head, next) {
        break;
      }
    }
  }
  free(free_node);
  return val;
}
```

Approach: fix the tail, then dequeue

# Lock-Free Queue Linearizability



- Still linearizable
- `dequeue()` "happens" with `atomic head = head->next`
- `enqueue()` "happens" with `atomic tail->next = node`
- No transient state is misleading, even the misplaced tail!
  - `dequeue()` and `enqueue()` both handle it




# The ABA Problem



# The ABA Problem

1. dequeue() wants to dequeue a node




```
elem* node = next;  
if (CAS(head, cur_head, next) {  
    return node;  
}
```



# The ABA Problem

2. Another dequeue() interrupts and dequeues A




Freed Nodes



# The ABA Problem

3. Another dequeue() interrupts and dequeues B




Freed Nodes




# The ABA Problem

4. An enqueue() re-uses A's memory for a new node
5. [Optional: other enqueue()s/dequeue()s occur]




Freed Nodes




# The ABA Problem

6. Finally, the original dequeue() completes




```
elem* node = next;  
if (CAS(head, cur_head, next) {  
    return node;  
}
```

Freed Nodes



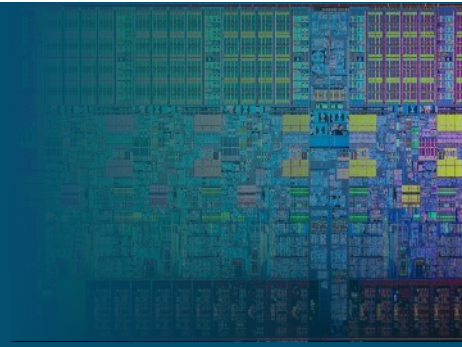
Still A, but next was B, and B was freed (and could contain anything!)

# The ABA Problem



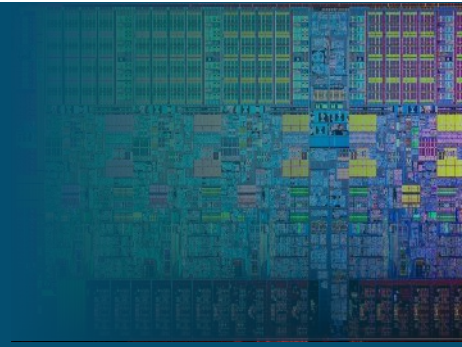
- Pointers are spatially unique, **not** also temporally unique
  - Re-using memory is necessary, but can cause unexpected problems
- If we re-use memory, we want the CAS to fail, even if the pointer points to the same memory
  1. Additional pointer information
  2. Free-list tracking

# The ABA Solution



- Version our pointers
  - 128-bit pointer
    - 64 bits of target address
    - 64-bit counter
  - Requires 128-bit CAS (64-bit x86 supports this)
- Track “freed” pointers in free list with associated counter
  - Re-use freed pointers but increment counter

# Lab 2 Check-In



- Lab 2: Performance testing
  - httpperf: <http://www.labs.hpe.com/research/linux/httpperf/httpperf-man-0.9.pdf>
  - Must run at most one httpperf client per machine
  - Try running on a different machine than your server
  - Experiment with parameters
    - We will provide a script to generate workloads



# Conclusions

- Turning a single lock into multiple (smaller-scope) locks can be done
  - Careful invariant consideration
  - Linearizability checking
- Removing locks entirely is possible
  - Atomicity still needed: CAS, TAS
  - Substantial engineering effort
- Next time:
  - Lock-free ordered lists
  - Lock-free hash tables