




CSCI-GA.3033-017  
**Special Topics:**  
**Multicore Programming**

**Lecture 8**  
**Synchronized Structures Part 2**

Christopher Mitchell, Ph.D.  
cmitchell@cs.nyu.edu || <http://z80.me>

# Homework 2 Review

## Thread-Safe Aggregation



```
1    static int sum_stat_a = 0;
2    static int sum_stat_b = 0;
3    static int sum_stat_c = 0;
4    int aggregateStats(int stat_a, int stat_b, int stat_c) {
5        sum_stat_a += stat_a;
6        sum_stat_b += stat_b;
7        sum_stat_c -= stat_c;
8        return sum_stat_a + sum_stat_b + sum_stat_c;
9    }
10   void init(void) { }
```

1. Use a single pthread mutex to make this function thread-safe. Add global variables and content to the `init()` function as necessary.
2. Modify the original code from to make it thread-safe, but use two mutices this time, one each for `sum_stat_a`, `sum_stat_b`, and `sum_stat_c`.

# Homework 2 Review

## Aggregation (1 Lock)




```
pthread_mutex_t stats_mutex;
1 static double sum_stat_a = 0;
2 static double sum_stat_b = 0;
3 static double sum_stat_c = 10000;
4 int aggregateStats(double stat_a, double stat_b, double stat_c) {
    pthread_mutex_lock(&stats_mutex);
5     sum_stat_a += stat_a;
6     sum_stat_b += stat_b;
7     sum_stat_c -= stat_c;
8     int rval = sum_stat_a + sum_stat_b + sum_stat_c;
    pthread_mutex_unlock(&stats_mutex);
    return rval;
9 }
10 void init(void) {
    pthread_mutex_init(&stats_mutex, NULL);
}
```

Invariant: The returned value from `aggregateStats()` must be equal to exactly the sum of the statistics at exactly the moment in time when the new additions were aggregated into each sum.



# Homework 2 Review

## Aggregation (3 Locks)



Invariant: The returned value from `aggregateStats()` must be equal to exactly the sums of statistic A, B, and C, each taken at the moment in time when the new addition was aggregated into **each** sum. The aggregate sum may therefore represent a value covering three close (but different) time periods.

```
sum_stat_a: stat_a_mutex  
sum_stat_b: stat_b_mutex  
sum_stat_c: stat_c_mutex
```

# Homework 2 Review

## Aggregation (3 Locks)



```
pthread_mutex_t stat_a_mutex;
pthread_mutex_t stat_b_mutex;
pthread_mutex_t stat_c_mutex;
static int sum_stat_a = 0;
static int sum_stat_b = 0;
static double sum_stat_c = 10000;
```

```
void init(void) {
    pthread_mutex_init(
        &stat_a_mutex, NULL);
    pthread_mutex_init(
        &stat_b_mutex, NULL);
    pthread_mutex_init(
        &stat_c_mutex, NULL);
}
```


```
int aggregateStats(
    int stat_a, int stat_b, int stat_c)
{
    int rval = 0;
    pthread_mutex_lock(&stat_a_mutex);
    sum_stat_a += stat_a;
    rval += sum_stat_a;
    pthread_mutex_unlock(&stat_a_mutex);
    pthread_mutex_lock(&stat_b_mutex);
    sum_stat_b += stat_b;
    rval += sum_stat_b;
    pthread_mutex_unlock(&stat_b_mutex);
    pthread_mutex_lock(&stat_c_mutex);
    sum_stat_c -= stat_c;
    rval += sum_stat_c;
    pthread_mutex_unlock(&stat_c_mutex);
    return rval;
}
```

# Outline

- volatile & Thread-Safe C++ Code
- A Lock-Free Hash Table



# Interlude: `volatile` and Thread-Safe C++ Code



- What does `volatile` mean in C?
  - Don't assume this value can't be changed elsewhere
- `volatile` means the same thing in C++
  - Rabid arguments about using `volatile` in C++  
eg: <http://stackoverflow.com/questions/4557979>
  - No memory fences, no order of execution
  - Some hardware may guarantee *tearing* does not occur
    - X86: Cache coherent
- Therefore: `volatile` useful only for rare single-word, many-write, "last to write wins" scenario w/ many cores





# Building Block: Concurrent Ordered List



# Concurrent Ordered Lists



- Idea: Build more complex data structures using thread-safe ordered lists
- Want a finer grained lock than the whole list
  - What's so hard about concurrent operations over lists?
- Let's explore:
  1. A lock-based approach
  2. A lock-free approach, based on our queue



# List-Based Set



- List is ordered
  - Inserts between any nodes
- Interface (looks roughly familiar?)
  - `bool insert(key): true if key didn't exist yet`
  - `bool lookup(key): true if key exists already`
  - `bool remove(key): true if key existed and was erased`
- Each function needs to locate key first
  - Proposed:  
`(node* prev, node* cur, node* next, bool found) lookup(key)`

# Lookup for Insertion



- Only prev's next really changes.
- Note: "head" is on the left, to make it clearer that this sorted list stores strictly increasing keys left-to-right






# Lookup for Removal

- Again, only prev's next really changes
- Tempting: only lock prev for update



# Hazard with Locks

- Concurrent deletion of 'B' and insertion of 'C'
- Naïve approach: Incorrect result!
- Hazardous deletion/deletion hazard as well



# Lock Coupling

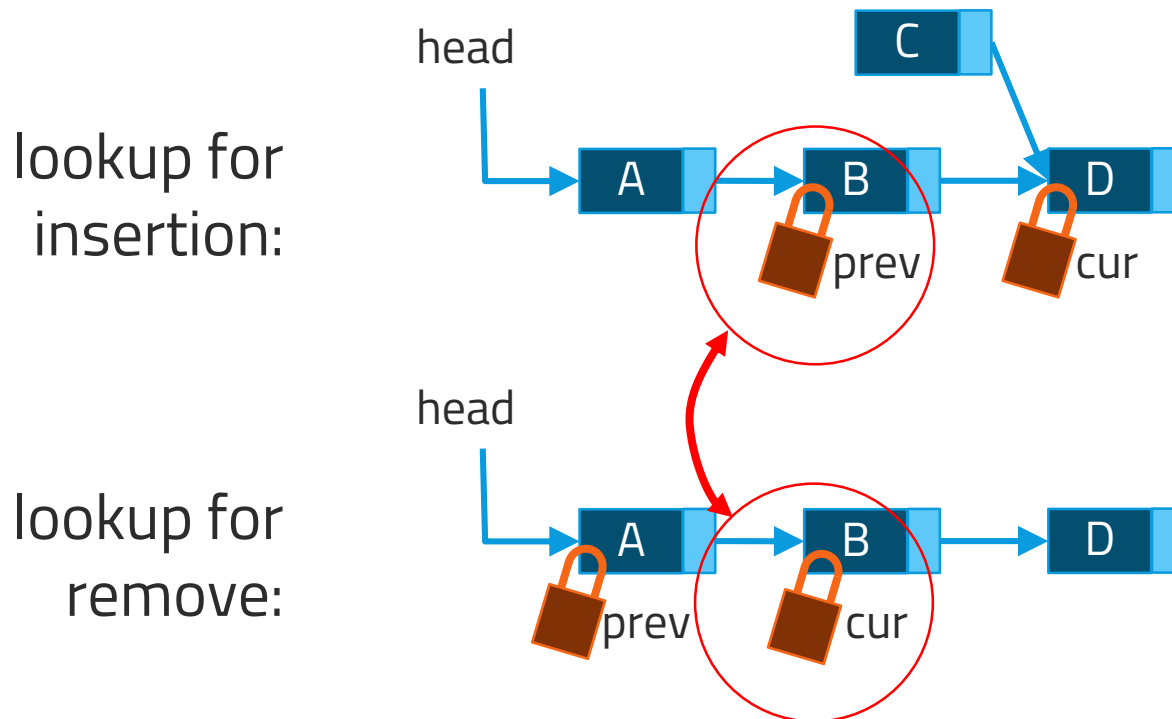


- Need to lock all nodes involved
  - On insertion, new node placed between prev and cur
  - On remove, prev node's pointer changes; cur disappears
- Acquire prev lock, then cur lock, during lookup.
- Lock on lookup, not the actual insertion/deletion!
- Still better than locking entire list




# Hazard Removed with Lock Coupling

- Concurrent deletion of 'B' and insertion of 'C'
- Lock-free version with CAS?




# Lock-Free Insertion



```
while(true) {  
    if find(list, key) {  
        return false // sets prev,cur  
    }  
    new_node->next = cur;  
    if CAS(&prev->next, cur, new_node) {  
        return true  
    }  
}
```

# Logical Deletion

1. Mark node deleted
2. Fix previous pointer



After step 1, want to make it that C can't be inserted after B since it's marked deleted. How?

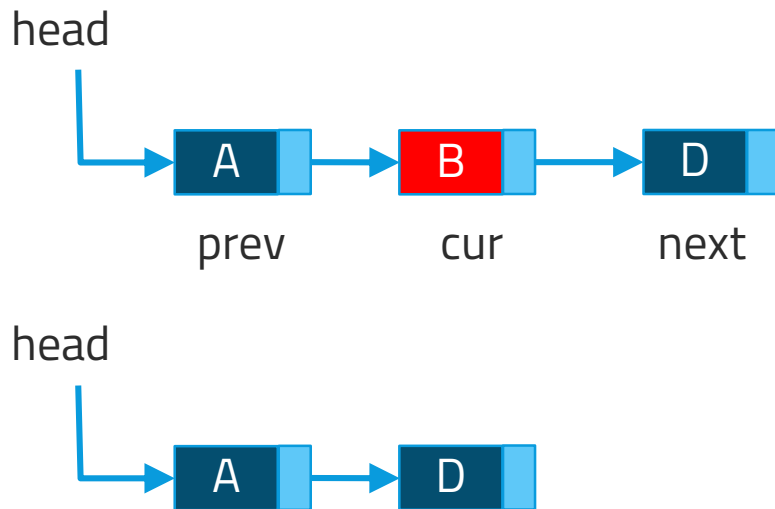


# Deletion Mark



- Combine next pointer and deletion mark (bit)
- All pointers are aligned!
  - Steal up to last 2 bits (for 32-bit pointers) or 3 bits (for 64-bit pointers)
- If the deletion bit is set a CAS against the original pointer (with low bit(s) reset) will fail
  - Works well for insert-delete hazards


# Lock-Free Deletion



```
while(1) {  
    if (!find(list, key)) {  
        // sets prev,cur,next  
        return false  
    }  
    if (CAS(&cur.mark, 0, 1)) {  
        break;  
    }  
}  
if (CAS(&prev->next, cur, next)) {  
    dealloc(cur);  
}  
return true
```

- Decouples logical and physical deletions via marker hidden in pointer
- Future find()s may see logically deleted nodes!

# Lock-Free Ordered List Summary



- Now have a basic building block for more complex structures
- Memory management notes
  - ABA problem recurs
  - Same solution: pointer-counter pairs






# Concurrent Hash Tables

# Concurrent (Locked) Hashing

- Closed-addressing hashing schemes
  - Put colliding elements in the same bucket
  - May need to change bucket count as more elements are hashed
- Open-addressing works too, but beyond the scope today




# Hash Table Lock Striping



- Approach: lock striping
  - Good results in practice
- Design pointers
  - Successful hash schemes maintain small buckets
  - No operation on the hash table affects all buckets
    - No *normal* operation (resizing...)
  - Each bucket could have its own lock (granularity)
  - Lookup needs locking! (why?)

# Hash Table: Array of Lock-Free Lists



- Use lock-free list as a bucket
- Same-bucket concurrent operations allowed
- Lookup is lock-free because of atomic CAS use
  - But not completely "free"

# Hashing and Resizing



- With lock striping
  - Resizing changes the number of buckets so it requires rehashing of all members
  - Grab all locks for resizing
  - Good results in practice but we want to use our lock-free lists
- Lock-free list
  - Can't grab locks to resize




# Hashing with Modulo

- First, hash keys using a good hashing algorithms
- Key hashes that share the same last  $\log_2 K$  bits have the same value modulo  $K$ 
  - Eg, let  $K = 4$  ( $\log_2 K = 2$ ).
  - $2 \% 4 = 30 \% 4 = 6 \% 4 = 2$  (binary 10)
  - $3 \% 4 = 11 \% 4 = 3$  (binary 11)
- Ordering keys within a bucket
  - Reverse the bits!
  - Eg, bucket 0b10 with  $K=4$ , holding 2, 30, and 6
  - 01000 (2), 01100 (6), 01111 (30)


# Hashing with Modulo: Split Ordered List

- Combine buckets into single ordered list
- Eg, buckets 0b10 and 0b11 for K=4
  - 0b10: 01000 (2), 01100 (6), 01111 (30)
  - 0b11: 11000 (3), 11010 (11)




# Resizing Split Ordered List

- Buckets can split, but the list stays the same
- Eg:  $K=4 \rightarrow K=8$ : Use last  $\log_2 8 = 3$  bits



# Concurrency During Resize

- Add additional layer(s) of indirection
- Skiplist



# Conclusion

- Concurrent ordered lists can be used for hashing and searching structures
- The ordered lists we saw today:
  - Decoupled deletion into two operations for lock-free operations
  - Allowed faster lookup but required locks for other operations
- Concurrent hash table
  - Different concurrency schemes can be combined in one structure
  - Difficult to show correctness