

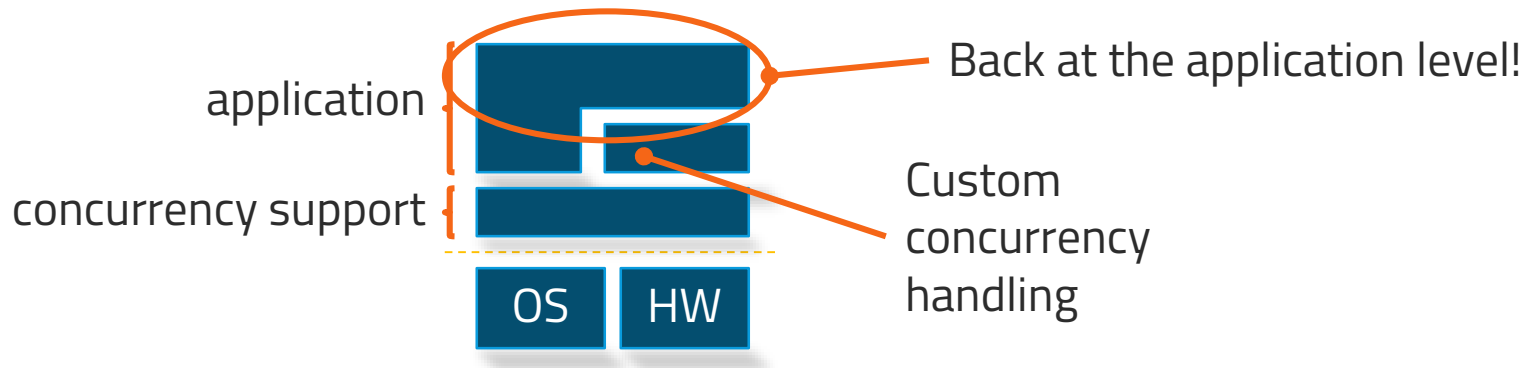


CSCI-GA.3033-017  
**Special Topics:**  
**Multicore Programming**

**Lecture 9**  
**Multicore Correctness**


Christopher Mitchell, Ph.D.  
cmitchell@cs.nyu.edu || <http://z80.me>

# Context



# Outline

- Taxonomy of Real-World Bugs
- Detecting and Reproducing Bugs
- Advanced Thread Interleaving
- Eliminating Non-Determinism?









# A Taxonomy of Real-World Bugs

# Types of Bugs



- Race conditions
- Deadlocks
- Atomicity Violations
- Ordering Violations
- Group Coordination Violations
- Timing Dependencies

# The Problem with Threads



"[Threads] discard the most essential and appealing properties of sequential computation: understandability, predictability, and determinism. Threads, as a model of computation, are wildly non-deterministic, and the job of the programming becomes one of pruning non-determinism"

-- Edward A. Lee, "The Problem with Threads", 2006

# A Simple Race Condition

- Accessing a shared variable outside a lock

```
Thread 1
{
    std::scoped_lock(m);
    i++;
}
```

```
Thread 2
    //...
    i--;
    //...
```

# A Simple Deadlock

- Inconsistent lock ordering

Thread 1

```
m1.lock();  
m2.lock();  
i++;  
m2.unlock();  
m1.unlock();
```

Thread 2

```
m2.lock();  
m1.lock();  
i--;  
m1.unlock();  
m2.unlock();
```



# More Subtle Bugs



- Some thread interleavings...
  1. Break implicit atomicity assumptions
  2. Break implicit order assumptions
  3. Break time interval guarantees
- Not every problem can be fixed with locks

# Atomicity Violation

- Find more examples in “Learning from Mistakes – A Comprehensive Study on Real-World Concurrency Bug Characteristics” (Lu, Park, Seo, Zhou 2009)

Thread 1

```
if (thr->proc_info) {  
    fputs(thr->proc_info);  
}
```

Thread 2

```
//...  
thr->proc_info = nullptr;  
//...
```

# Atomicity Violation

- Find more examples in “Learning from Mistakes – A Comprehensive Study on Real-World Concurrency Bug Characteristics” (Lu, Park, Seo, Zhou 2009)

Thread 1

```
1 if (thr->proc_info) {  
2   fputs(thr->proc_info);  
}
```

Thread 2

```
//...  
3 thr->proc_info = nullptr;  
//...
```

- Assumes 3 will not be interposed between 1 and 2

# Order Violation

- Implicit programmer expectations

Thread 1

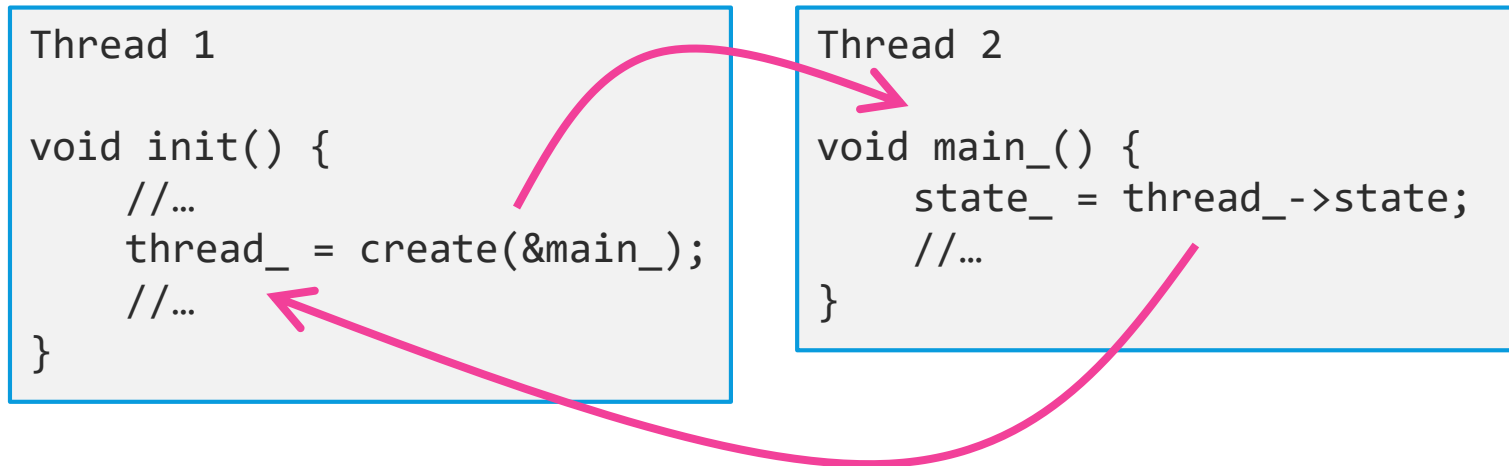
```
void init() {  
    //...  
    thread_ = create(&main_);  
    //...  
}
```

Thread 2

```
void main_() {  
    state_ = thread_->state;  
    //...  
}
```

# Order Violation

- Implicit programmer expectations



- What if `create()` doesn't return until `main_` runs for a while?



# Another Order Violation

- Subtle write-write race

Thread 1

```
int readWriteProc() {  
    //...  
    ReadAsync(&p);  
    io_pending = true;  
    while(io_pending) {  
        // Wait for done  
        //...  
    }  
}
```

Thread 2

```
void doneWaiting();  
    // Callback called from  
    // ReadAsync()  
    io_pending = false;  
}
```

# Another Order Violation

- Subtle write-write race

Thread 1

```
int readWriteProc() {  
    //...  
    ReadAsync(&p);  
    1 io_pending = true;  
    while(io_pending) {  
        // Wait 2 for done  
        //...  
    }  
}
```

Thread 2

```
void doneWaiting();  
    // Callback called from  
    // ReadAsync()  
    3 io_pending = false;  
}
```

- Programmer assumes 1 must run before Thread 2 manages to get to 3

# Group Coordination Bugs



Thread 1

```
void destroyCtx() {  
    references--;  
    if (!references) {  
        free(&resource);  
    }  
}
```

Thread 2

```
void destroyCtx() {  
    references--;  
    if (!references) {  
        free(&resource);  
    }  
}
```

# Group Coordination Bugs



Thread 1


```
void destroyCtx() {  
  2 references--;  
  3 if (!references) {  
    free(&resource);  
  }  
}
```

Thread 2

```
void destroyCtx() {  
  1 references--;  
  4 if (!references) {  
    free(&resource);  
  }  
}
```

- Imagine ordering of 1, 2, 3, 4: both Thread 1 and Thread 2 will try to free the resource
- Type of race condition

# Timing Dependencies



- Many threads may cause timeout to spuriously trigger

Thread i

```
//...  
rw_lock(m);  
//...
```

Thread i

```
//...  
{  
    try_lock_for(m);  
}  
  
// timeout
```



# Timing Dependencies

- Many threads may cause timeout to spuriously trigger

Thread i

```
//...  
rw_lock(m);  
//...
```

Thread i

```
//...  
rw_lock(m);  
//...
```

Thread i

```
//...  
{  
    try_lock_for(m);  
}  
  
// timeout
```



# Detecting and Reproducing Bugs

# Finding a Bug-Fix Strategy



- First, we must detect bugs. How?

“Three quarters (73%) of the examined non-deadlock bugs are fixed by techniques other than adding/changing locks. Programmers need to consider correctness, performance and other issues to decide the most appropriate fix strategy.”

-- “Learning from Mistakes – A Comprehensive Study on Real-World Concurrency Bug Characteristics” (Lu, Park, Seo, Zhou 2009)

# Software-Based Detection

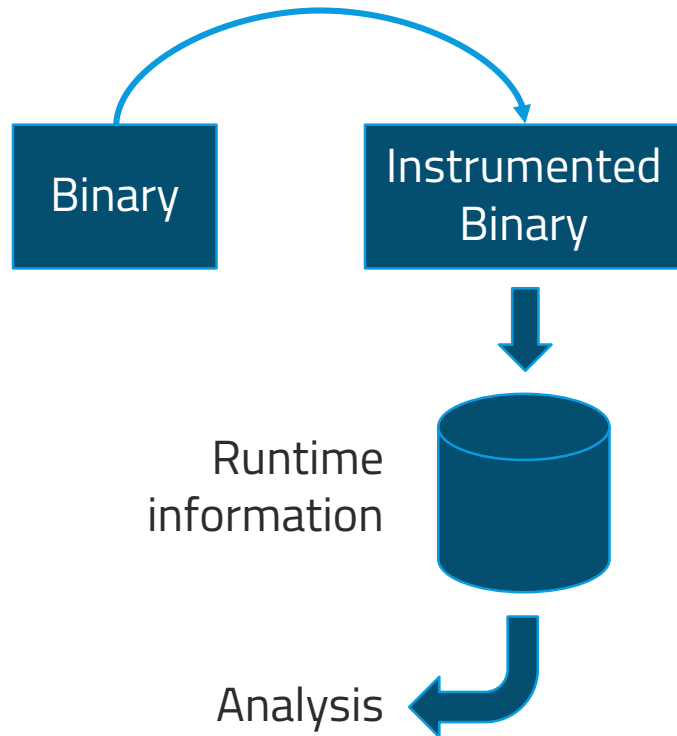


- Static analysis: inspect during compilation
- Dynamic analysis: inspect during runtime
  - Catches more than static checking
  - Shared variables may not always be static (e.g.: pointers)
  - Subtleties of shared variable protection that cannot be captured by static analysis
    - Anything involving non-deterministic input
- Binary instrumentation
- Dynamic binary translation

# Software-Based Detection




- Binary Instrumentation




- Eg: helgrind

- Dynamic binary translation

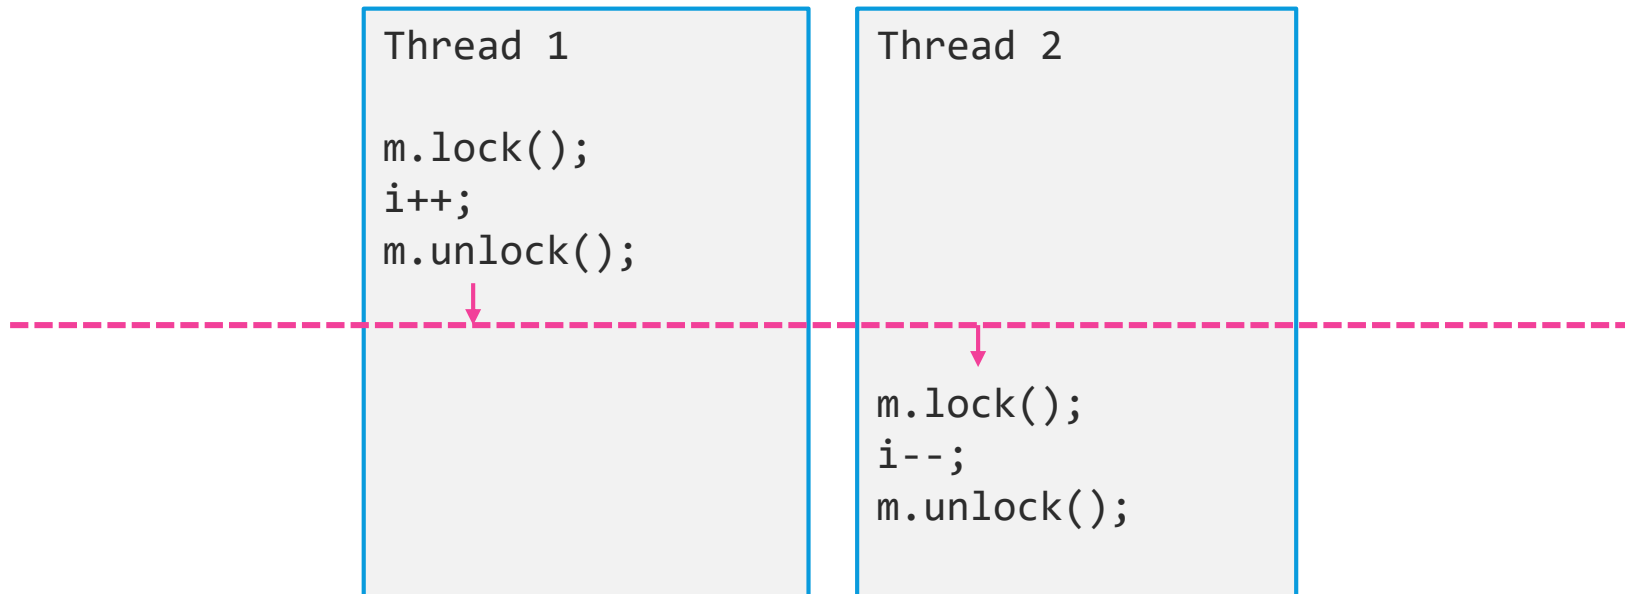




# "Happens-Before" Graphs

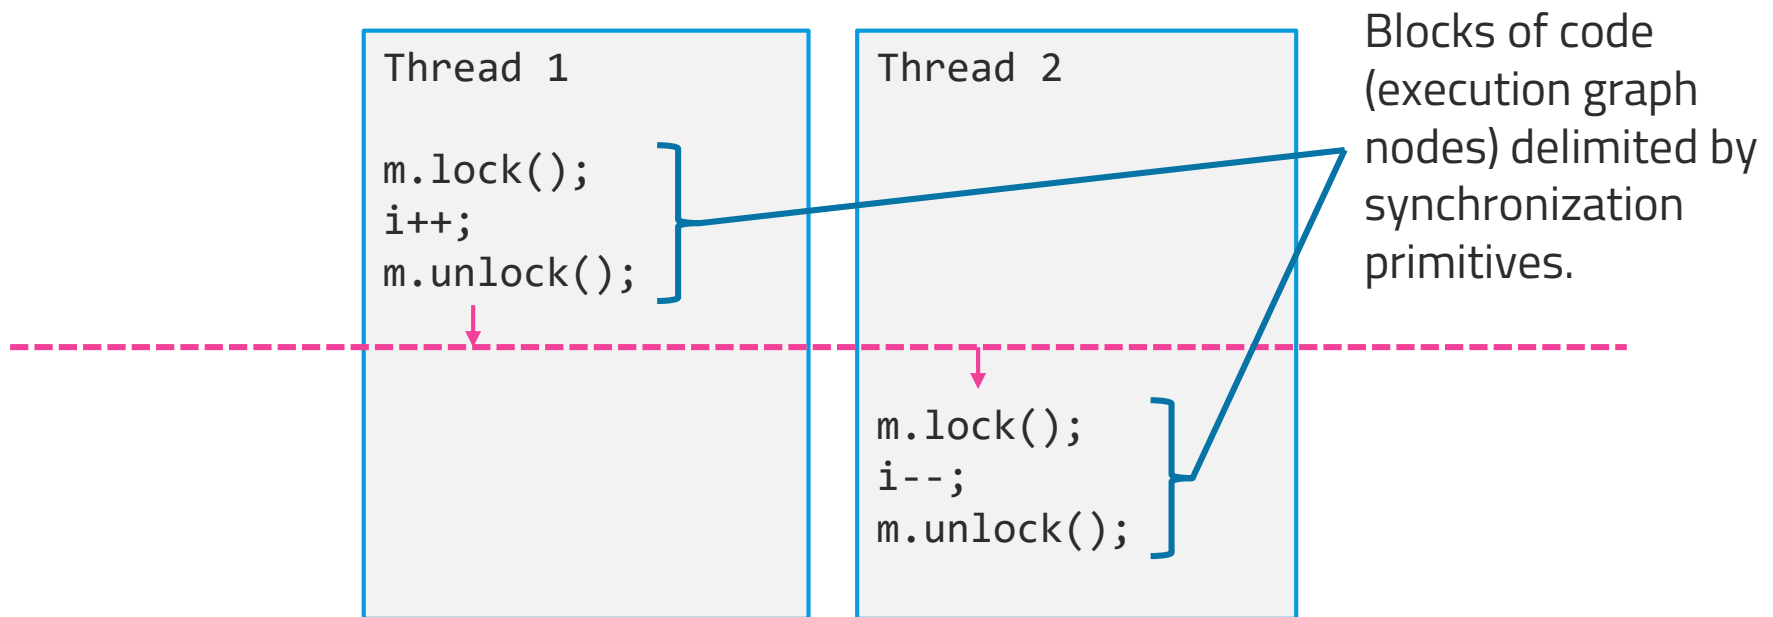


- Basic concept also used in software analysis



# "Happens-Before" Graphs

- Basic concept also used in software analysis



# Race Detection



- `i` accessed by both nodes, but they do not have a “happens before” relationship

Thread 1

```
{  
    scoped_lock(m);  
    i++;  
}
```

Thread 2

```
//...  
i--;  
//...
```

- Other examples: “Eraser: A Dynamic Data Race Detector for Multithreaded Programs”, Savage et al., 1997.

# A More Subtle Hazard



Thread 1

```
count++;  
m.lock();  
i++;  
m.unlock();
```

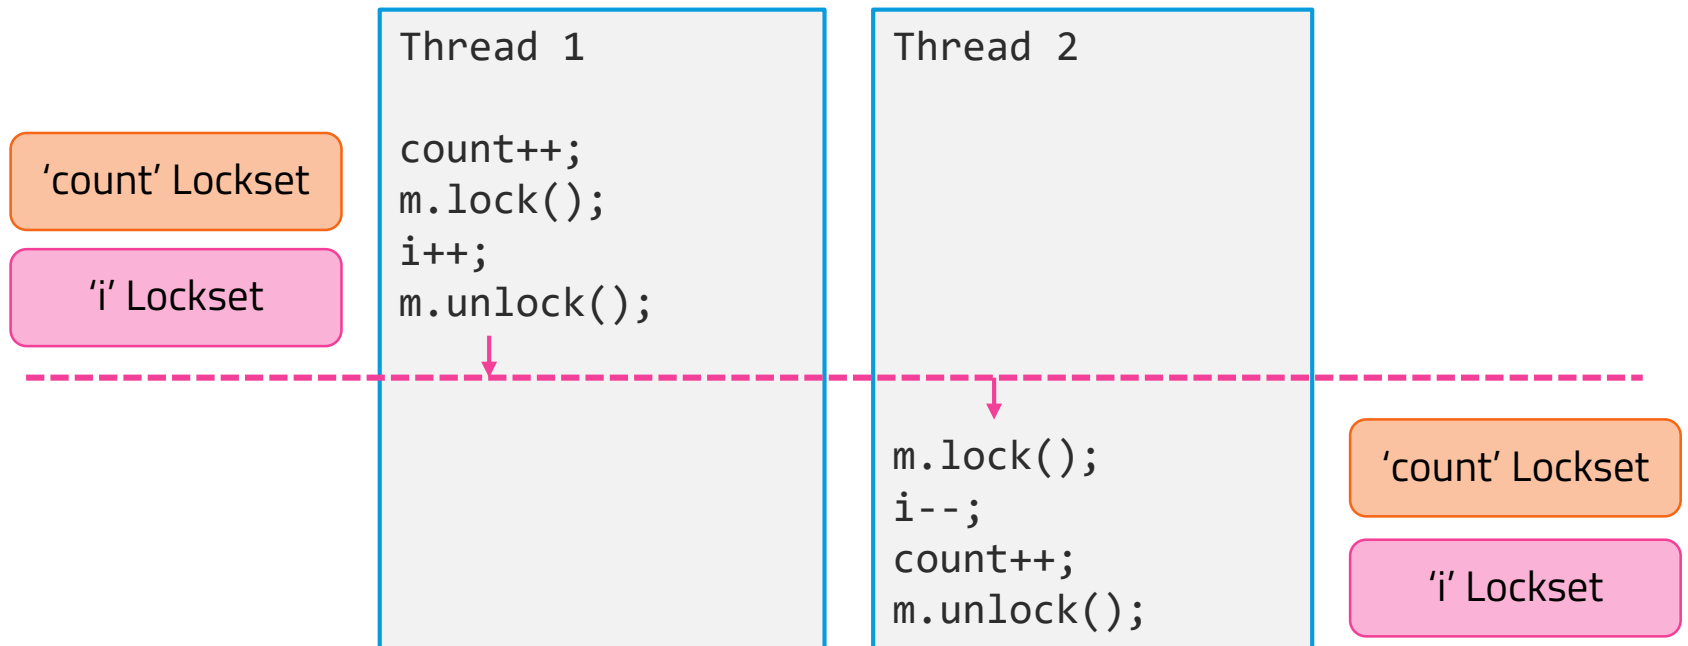
Thread 2

```
m.lock();  
i--;  
count++;  
m.unlock();
```



# Locksets

- A record for each variable read/written
  - Performed under which lock(s)?
  - Performed in which block?
- Managing lockset size





# Further Detection Needed



- How to know 3 is at fault for a crash at 2?

Thread 1

```
1 if (thr->proc_info) {  
2   fputs(thr->proc_info);  
}
```

Thread 2

```
//...  
3 thr->proc_info = nullptr;  
//...
```

- Locksets can be refined to capture situations that are races, but are harmless
- Some bugs involve atomicity violations or order violations, not incorrect lock use

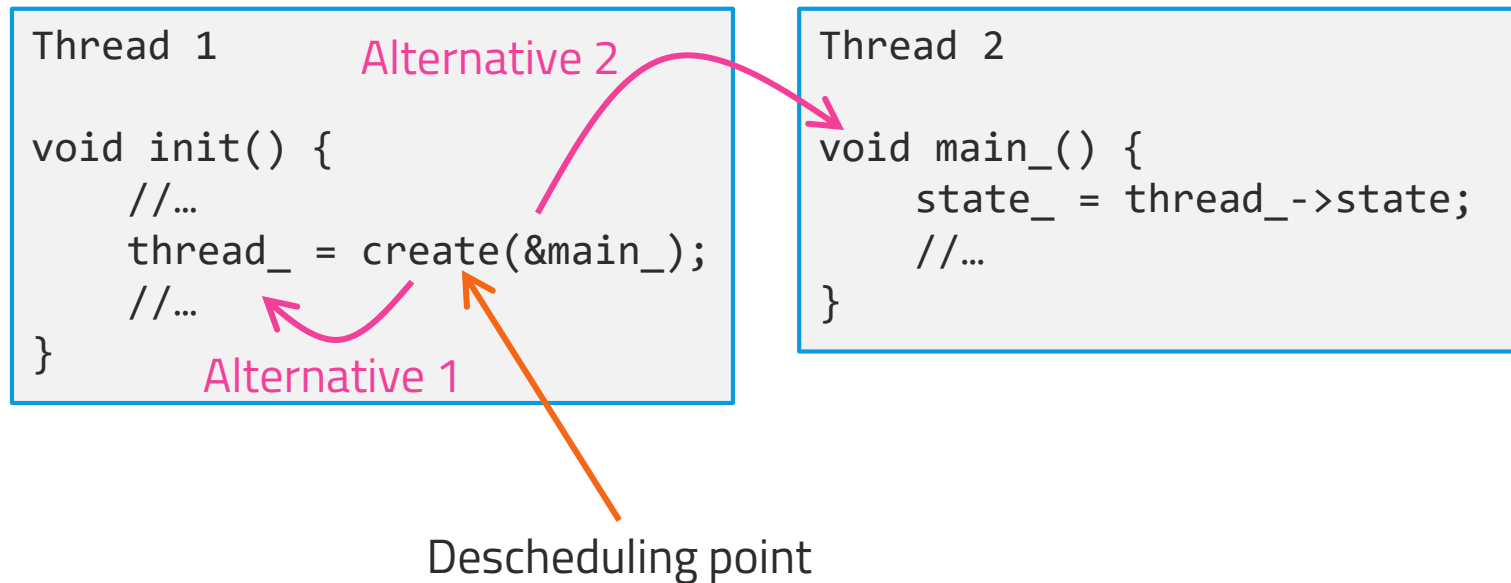
# Record/Replay



- Trace backwards from where bug manifested to track root cause (eg, what set `thd->proc_info` to `nullptr`)
- Requires storing huge amounts of state


# A Software Approach

- Using threading API (e.g., pthreads) to explore potential hazards



# Scheduler-Based Framework

- See Musuvathi et al, "Finding and Reproducing Heisenbugs", 2008.




# Taming State Explosion




- A program with  $n$  threads that execute  $k$  atomic steps has  $n^k$  possible interleavings
- If we reduce the number of preemptions,  $k$  decreases sharply.
  - Tradeoff of coverage and analysis time
- Empirical evidence: few threads necessary to expose atomic and order violations.

# Simplifying Parallel Programming




- Two of many efforts
  - Hardware Transactional Memory (future lecture):  
Research dates back to 1993
  - Deterministic execution: guaranteeing deterministic semantics in parallel software
- More efforts: to be mentioned in your presentations!

# Deterministic Execution



- Recently, arguments for exploring deterministic ways to express parallelism
  - “Parallel Programming Must Be Deterministic By Default”, Bocchino et al, 2009.
- Language itself would have constructs for compile-time enforcements of sharing constraints
- Ongoing effort, with many recent publications

# Conclusion



- At some level (ideally as low as possible), threads must exist
  - Hardware primitive: multiple cores
- Continuous, wide effort to expose different model to higher-level programmer
  - Programmer still wants parallel view of the world
- Main challenge: Taming non-determinism inherent in pure thread model



# References



- Edward A. Lee, "The Problem with Threads", 2006
- Lu, Park, Seo, Zhou, "Learning from Mistakes – A Comprehensive Study on Real-World Concurrency Bug Characteristics", 2009
- Savage et al, "Eraser: A Dynamic Data Race Detector for Multithreaded Programs", 1997
- Musuvathi et al, "Finding and Reproducing Heisenbugs", 2008.
- Bocchino Jr. et al, "Parallel Programming Must Be Deterministic By Default", 2009.
- Helgrind: <http://valgrind.org/docs/manual/hg-manual.html>