



CSCI-GA.3033-017  
**Special Topics:**  
**Multicore Programming**

**Lecture 12**  
**Transactional Memory**

Christopher Mitchell, Ph.D.  
cmitchell@cs.nyu.edu || <http://z80.me>

# Database Background



- Databases have successfully exploited parallel hardware for decades
  - Multi-core computers
  - Multi-machine clusters
- Database performance optimizations
  - Execute queries concurrently
  - Run queries on multiple CPUs
- Query author need not know about parallelism

# Database Background



- Database programming model: transactions
  - Computation executes as if it was the only computation accessing the database
- Query results are indistinguishable from serial execution
  - Serializability!
- Transactions allow concurrent operations to access a common database and still produce predictable, reproducible results



# Transactional Memory



## Transactions

- Inspired by databases
- Coordinate concurrent access
- Characteristics: ACID
  - Atomicity
  - Consistency
  - Isolation
  - Durability

## Memory

- Main data storage during typical multicore programs
- Computations wrapped in transactions
- Much more primitive than database transactions

# Transactions: ACID

A transaction is a sequence of actions that appears indivisible and instantaneous to an outside observer.

## Atomicity

All constituent actions in a transaction complete successfully, or none appears to start executing.

## Consistency

(Semantics depend on the application)

## Isolation

Transactions do not interfere with each other while they are running, regardless of whether they are operating in parallel or not.

## Durability

Once a transaction commits, its result is permanent.

# Transaction Example #1

## Thread 1 (Query 1)

```
transaction.begin();  
a = a - 20;  
b = b + 20;  
a = a - b;  
c = c + 20;  
transaction.end();
```

Thread 1's (Query 1's) accesses and updates to a, b, and c are atomic.

## Thread 2 (Query 2)

```
transaction.begin();  
c = c - 30;  
a = a + 30;  
transaction.end();
```

Thread 2 (Query 2) sees either all or none of Thread 1's (Query 1's) updates.

# Transaction Example #2

```
int x = 0, y = 0;
```

## Thread 1 (Query 1)

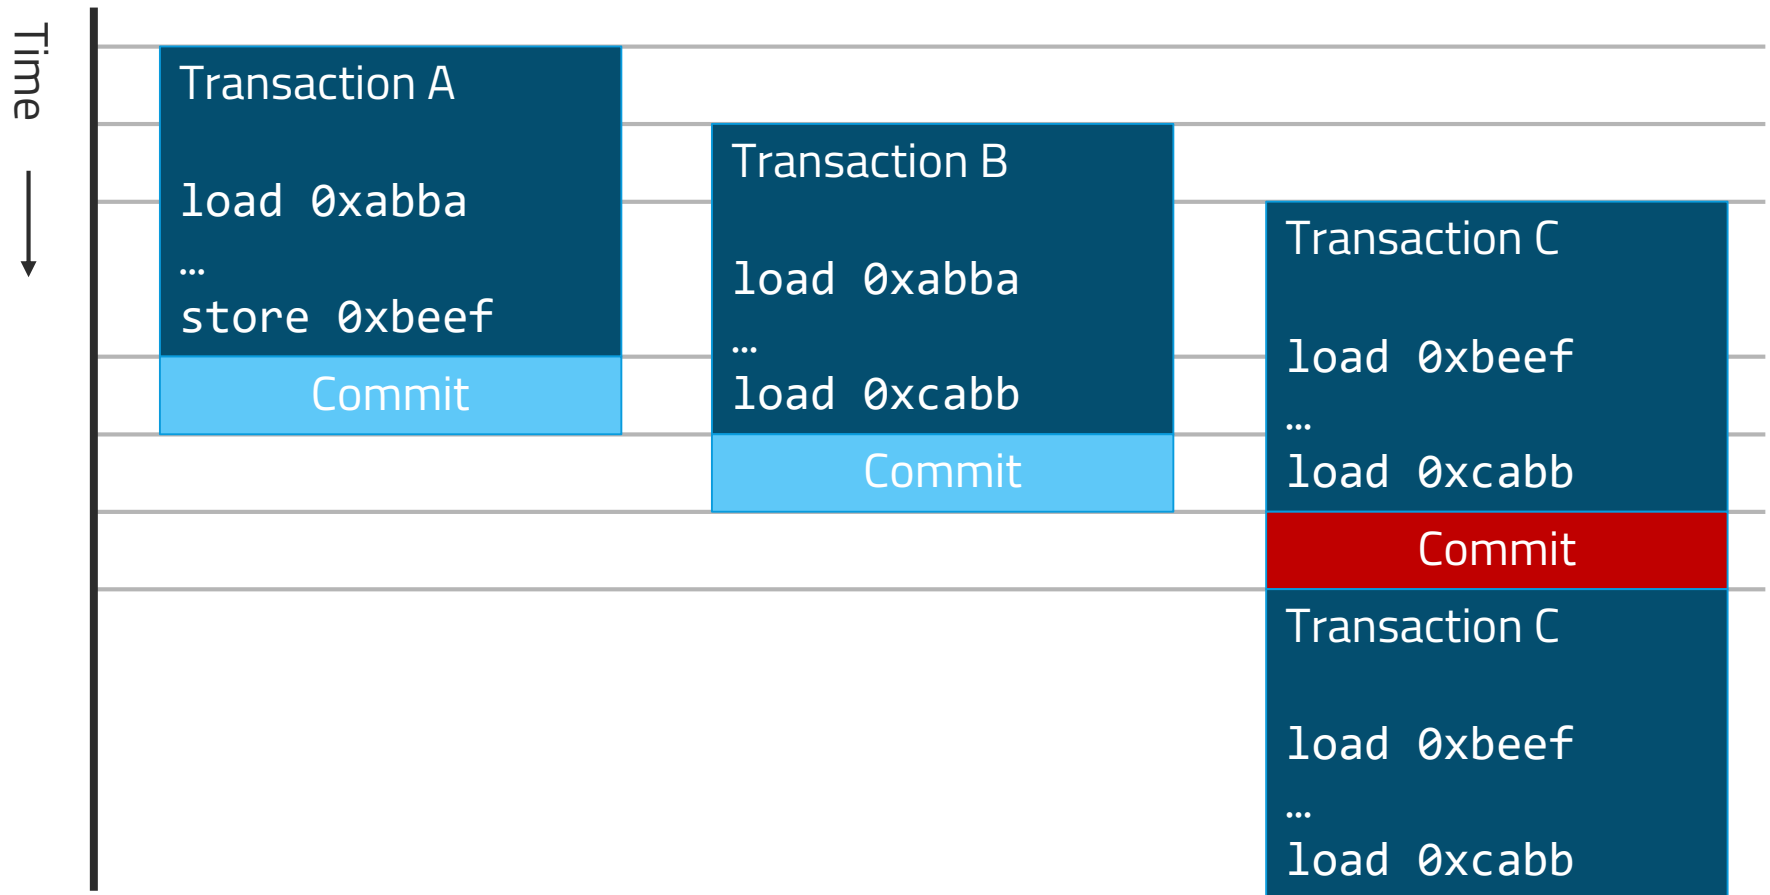
```
transaction.begin();  
x = 42;  
y = 42;  
transaction.end();
```

## Thread 2 (Query 2)

```
transaction.begin();  
tmp1 = x;  
tmp2 = y;  
transaction.end();
```

What values does Thread 2 (Query 2) see?

# Transaction Example #3





# Who Uses Transactional Memory?




- Programmer
- Compiler designer
  - To implement high-level language features
- Ancillary uses (general parallel programming)
  - Error recovery
  - Realtime programming
  - Multitasking

# Transactional Memory in GCC/G++



- Specification: <http://bit.ly/2gDWjR5>
- `__transaction_relaxed {  
    statements  
}`
- `__transaction_atomic {  
    statements  
}`
- `__transaction_cancel;`
- Works with both software and hardware transactional memory (more on that later)

# Atomic Transactions



- Other code cannot see intermediate results from atomic transaction.
- Do not see intermediate results from other code.
- Limited subset of transaction-safe operations
- Can be cancelled

# Relaxed Transactions



- Other code cannot see intermediate results from relaxed transaction.
- Do not see intermediate results from other code.
- Any operation allowed
- Cannot be cancelled
- Act as if all relaxed transactions use a single mutex (serialization)

# Transactional C/C++ Functions

```
__attribute__((transaction_callable)) double foo(void);

double bar(void) {
    ...
    __transaction_atomic {
        ...
        foo();
        ...
    }
    ...
}
```



# Concurrency Control: Conflict Detection/Resolution



- A conflict occurs when two transactions perform conflicting operations on the same piece of memory
  - 2 writes
  - 1 read, 1 write
- The conflict is detected by the underlying Transactional Memory (TM) system
- The conflict is resolved by the TM, e.g. by aborting, delaying, or repeating one or more transactions.

# Concurrency Control: Conflict Detection/Resolution



- Motivating example: double-ended queue



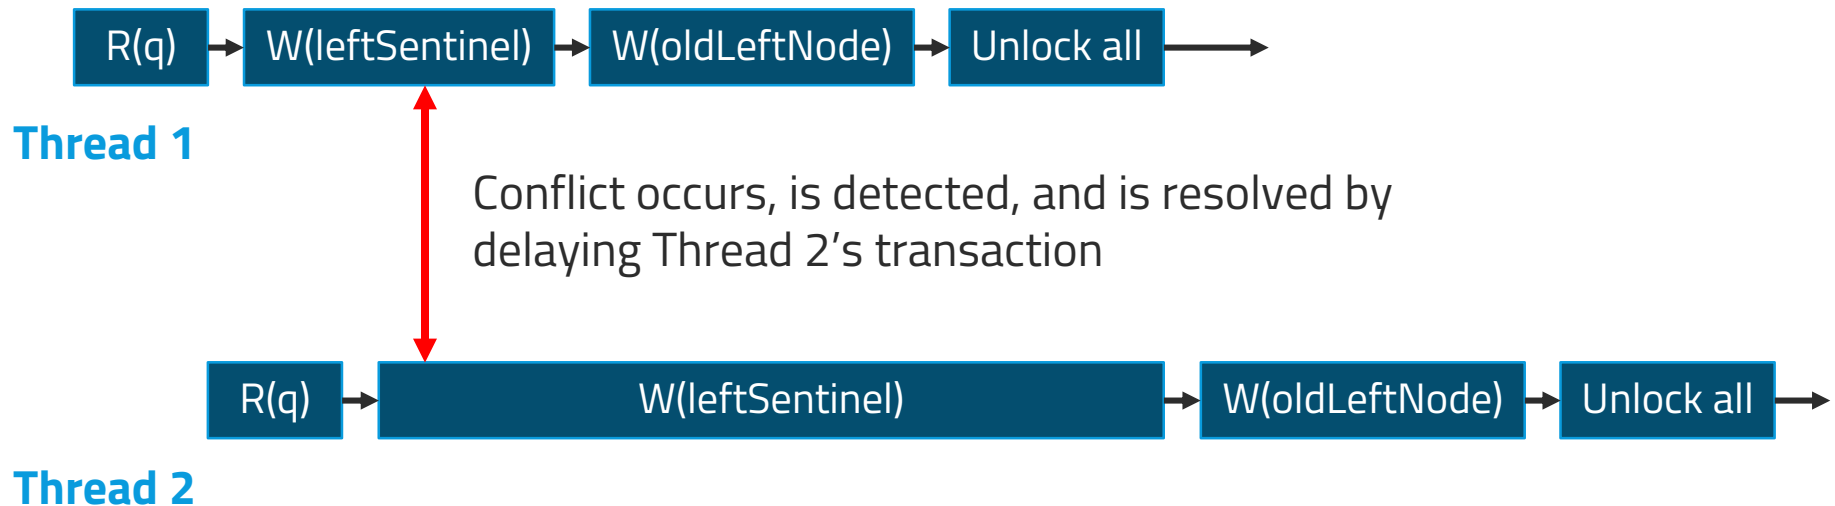
```
void pushLeft(Dqueue *q, int val) {
    QNode *qn = new Qnode;
    qn->val = val;
    QNode *leftSentinel = q->left;
    QNode *oldleftNode =
        leftSentinel->right;
    qn->left = leftSentinel;
    qn->right = oldleftNode;
    leftSentinel->right = qn;
    oldleftNode->left = qn;
}
```



```
void pushLeft(Dqueue *q, int val) {
    QNode *qn = new Qnode;
    qn->val = val;
    __transaction_atomic {
        QNode *leftSentinel = q->left;
        QNode *oldleftNode =
            leftSentinel->right;
        qn->left = leftSentinel;
        qn->right = oldleftNode;
        leftSentinel->right = qn;
        oldleftNode->left = qn;
    }
}
```

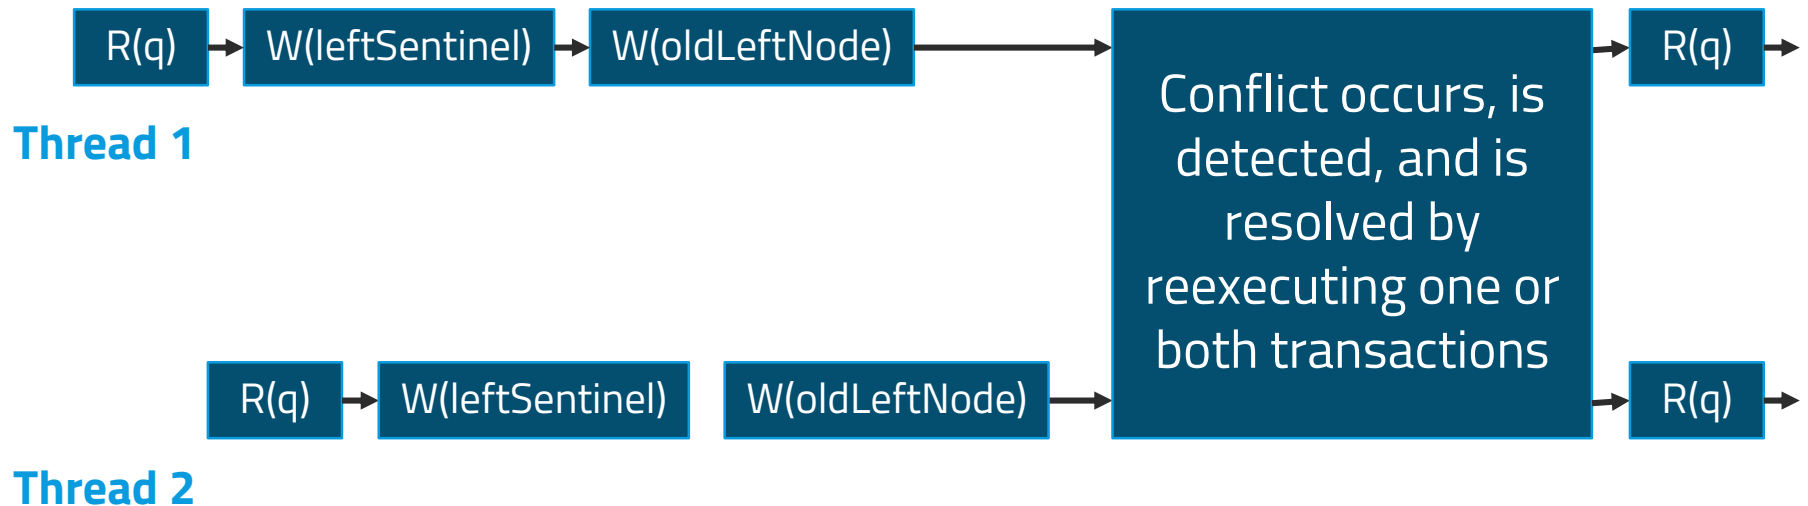
# Resolving Conflicts: Pessimistic Concurrency Control

- Some TMs use optimistic control, others use pessimistic control



# Resolving Conflicts: Optimistic Concurrency Control

- Some TMs use optimistic control, others use pessimistic control



# Classifying Available TMs




Read	Write ->	Optimistic	Pessimistic
Optimistic		TCC* TL2* SigTM*	Intel C++ STM Intel Java STM HASTM* Microsoft STM.NET
Pessimistic		LogTM*	Intel C++ STM

*\* = academic project or publication; unmarked items are public products*



# Version Management



- How to handle writes before a transaction commit?
- Eager version management
  - Transaction directly modifies memory
  - Maintains undo log of overwritten data
  - Requires pessimistic concurrency control
- Lazy version management
  - Transaction records redo log of writes to apply
  - Updates delayed until transaction commits
  - Works with optimistic concurrency control

# Conflict Detection




- Pessimistic approach: Easy
  - Locks
- Optimistic approach: Harder
  - Granularity of conflict (cache line, object, etc.)
  - Detection interval
    - When read/write set is declared: eager conflict detection
    - On validation: up to several instances throughout transaction
    - On commit: lazy conflict detection
  - Conflicting accesses
    - Concurrent transactions
    - Concurrent and committed transactions


# Providing TM



## Software Transactional Memory (STM)



## Hardware Transactional Memory (HTM)



# Software Transactional Memory (STM)



- Compiler instruments code with transaction prolog, epilog, and read/write functions
- Runtime tracks memory accesses, detects conflicts, and commits/aborts execution.

```
__transaction_atomic {  
    r = x;  
    y = r + 1;  
}
```



```
td = getTxnDesc();  
txnBegin(td);  
r = txnRead(td, &x);  
txnWrite(td, &y, r + 1);  
txnEnd(td);
```

# Software Transactional Memory (STM)



- Components:
  - Transaction descriptor: per-transaction data structure
  - Undo log (eager version management)
  - Redo log (lazy version management)
  - Readset/writeset: tracks memory locations the transaction has read or written.




# Hardware Transactional Memory (HTM)



- Three flavors
  1. Full hardware TM
  2. Software and hardware TM used together
  3. Hardware extensions to accelerate STM

# Hardware Transactional Memory (HTM) Functions



- Identify memory locations for transactional accesses
- Manage readsets and writesets of transactions
- Detect and resolve data conflicts
- Manage architectural register state
- Commit or abort transactions

# Supporting Transactions



Requirement	Satisfied By
Buffering	Transactional cache
Conflict Detection	Cache coherence protocol
Abort/Recovery	Invalidate transactional cache line
Commit	Validate transactional cache line

# Hardware Transactional Memory (HTM)



- Extensions to the instruction set
- Tracking readsets and writesets using caches and buffers
- Coherence messages trigger conflict detection
- Nearly all existing HTMs perform eager conflict detection

# IBM BlueGene/Q and HTM



- Multicore 64-bit PowerPC-based SoC
- 18 cores
  - 16 for computations
  - 1 for OS
  - 1 spare
- HTM in 32MB L2 cache



# Intel Haswell and HTM



- Intel Haswell (2013): First consumer CPU to support HTM
  - Later disabled in Haswell and early Broadwell due to bugs
- Transactional Synchronization Extensions (TSX)
  - Hardware Lock Elision (HLE): Backwards-compatible conversion from lock-based programs to transactional programs
  - Restricted Transactional Memory (RTM): Complete, non-backwards-compatible transactional memory.

# STM Advantages Over HTM



- More flexible than hardware
  - Permits implementation of a wider variety, more sophistication algorithms
- Easier to modify and evolve
- Integrate more easily with existing systems and language features, such as garbage collection
- Fewer intrinsic limitations imposed by fixed-size hardware structures, such as caches.

# HTM Advantages Over STM




- Execute applications with lower overhead
- Less reliant on compiler optimizations to provide performance
- Better power and energy profiles
- Can treat all memory accesses in a transaction as *implicitly* transactional
- Strong isolation without changing non-transactional memory accesses
- Well-suited to languages (like C/C++) without dynamic compilation, garbage collection, etc.

# Conclusion

- Transactional Memory: Strong candidate for parallel programming
- Many different implementations, poor portability
  - Still a work-in-progress; details vary between languages and processors
- Trades bandwidth for simplicity and latency tolerance
- Transactions eliminate locks
  - Transactions are inherently atomic
  - Catches most common parallel programming errors
- Shared memory consistency is simplified
- Shared memory coherence is simplified

# Reminders



- 12/11: *Optional*. Brief final exam review will be provided.
- 12/12: Presentations
- 12/18: Lab 4 due; final exam